

A JAX–PETSc Toolset for Nonlinear Finite-Element Energy Minimization:

Local Automatic Differentiation, Sparse Assembly, and Solver Policy

Michal Béréš

Institute of Geonics of the Czech Academy of Sciences, Ostrava, Czech Republic

Department of Applied Mathematics, Faculty of Electrical Engineering and Computer Science,

VSB – Technical University of Ostrava, Ostrava, Czech Republic

michal.beres@ugn.cas.cz

Abstract

Large nonlinear finite-element energy problems require reliable local derivatives, robust nonlinear globalization, and sparse distributed linear algebra that is matched to the mechanics of the problem. We develop a scientific toolset for studying these requirements as coupled numerical choices in nonlinear finite-element energy minimization and energy-based topology optimization. The primary JAX+PETSc realization uses JAX for local energy and constitutive kernels and PETSc for sparse assembly, Newton globalization, Krylov solvers, and preconditioners. The benchmark suite covers p -Laplace, Ginzburg–Landau, finite-strain hyperelasticity, two- and three-dimensional Mohr–Coulomb plasticity, and topology optimization. Where matched formulations are available, pure JAX and FEniCS provide internal reference comparisons. JAX-FEM provides the external hyperelastic comparison, while the Sysala-family slope-stability literature and the Čermák–Sysala–Valdman MATLAB work provide reference context for the plasticity endpoint observables.

Across these problems, we compare element automatic differentiation, constitutive automatic differentiation, and colored sparse recovery for constructing second-order information, and measure their interaction with PETSc assembly, nonlinear globalization, Krylov solvers, and preconditioners. In the tested three-dimensional plasticity case, constitutive AD has the lowest measured wall time while preserving the reported terminal observables. The distributed mechanics and topology examples provide MPI timing, solver-behavior, and rank-consistency evidence under the stated protocols.

1 Introduction

Nonlinear finite-element models in mechanics, scalar PDEs, and design optimization often combine three computational difficulties. The local model may be easiest to write as an energy or constitutive law whose derivatives are tedious to assemble by hand. The global problem then requires robust nonlinear globalization and sparse Newton–Krylov linear algebra. Finally, the implementation must remain comparable across reference formulations, derivative approximations, and distributed-memory solver choices. A useful scientific toolset for this class of problems should therefore expose not only a modeling interface, but also the derivative route, assembly layout, nonlinear solver policy, and preconditioner design.

We develop and evaluate such a scientific toolset for nonlinear finite-element energy minimization and optimization. In the primary JAX+PETSc realization, JAX defines local energies and constitutive laws, while PETSc provides sparse distributed algebra, nonlinear globalization, Krylov methods, and preconditioners. The pure JAX and FEniCS formulations provide internal checks where matched benchmark definitions exist. External comparisons are deliberately scoped: JAX-FEM is used for the hyperelastic benchmark, and the Sysala-family slope-stability literature, including the Čermák–Sysala–Valdman MATLAB work, provides context for the plasticity endpoint observables.

1.1 Contributions

Our contribution is a JAX+PETSc framework that treats derivative construction, MPI sparse assembly, nonlinear globalization, and preconditioning as coupled design choices for nonlinear FEM energy minimization. It is organized around three solver components:

1. **Globalized energy minimization.** We compare line-search, trust-region, and hybrid globalization choices on representative nonlinear finite-element benchmarks, with conclusions restricted to the tested benchmark definitions.
2. **Derivative construction in sparse assembly.** We implement and compare element automatic differentiation, constitutive automatic differentiation, and colored sparse finite-difference recovery inside the same sparse assembly path.

3. Distributed PETSc realization. We connect the nonlinear and derivative choices to PETSc Krylov methods and preconditioners, including Hypr, GAMG, and problem-specific polynomial multigrid policies used in the larger mechanics and topology runs.

This organization lets the numerical experiments vary derivative information and solver policy without changing the benchmark definition.

We evaluate these components on six benchmark families: p -Laplace, scalar Ginzburg–Landau, hyperelasticity, two- and three-dimensional Mohr–Coulomb-type plasticity, and topology optimization. The scalar problems isolate formulation and derivative agreement. Hyperelasticity adds finite-strain mechanics with a matched JAX-FEM comparison. The plasticity benchmarks test nonsmooth constitutive response, endpoint-observable comparison, and multigrid behavior. Topology optimization tests a distributed design-and-mechanics loop. The geomechanics context follows the finite-element strength-reduction and 3D slope-stability literature [1–5], but the present plasticity study is limited to final-load endpoint observables rather than path-consistent incremental plastic-history validation.

The surrounding literature covers expressive finite-element automation, differentiable FEM, JAX–solver bridges, mechanics and topology benchmark families, and sparse nonlinear solver infrastructure. The present work combines these roles by treating local differentiation, PETSc-owned sparse MPI assembly, nonlinear globalization, and preconditioner policy in one heterogeneous benchmark suite. Section 2 gives the source-specific discussion and records the role map used to position the surrounding software and literature.

The paper is organized as follows. Section 2 deepens the literature context behind this positioning. Section 3 describes the common finite-element energy notation, derivative routes, nonlinear globalization, and sparse derivative recovery. Section 4 explains the primary JAX+PETSc implementation and the narrower pure JAX and FEniCS reference implementations used in selected checks. Section 5 defines the benchmark suite, with each benchmark separated into literature model and implemented discrete problem. Section 6 then reports validation and external-comparison evidence, deliberately separating hyperelastic JAX-FEM agreement from the three-dimensional plasticity endpoint-surrogate comparison. Section 7 reports timing, scaling, derivative-route, and discretization-comparison evidence. Section 8 and Section 9 synthesize the main scope and limitations. Supporting solver-policy diagnostics are collected in Section A.

2 Related Work

The relevant literature falls into four groups: finite-element and PDE-optimization automation, differentiable FEM and JAX–solver bridges, mechanics and topology benchmark context, and sparse nonlinear solver infrastructure. We use this taxonomy to position the present toolset without ranking the surrounding software ecosystems.

Table 1 records the computational roles needed for this positioning: modeling layer, differentiation scope, second-order information, parallel execution model, and benchmark families closest to the present suite. The table is a role map rather than a package ranking; the paragraphs below give the source-specific discussion.

Technical family	Documented role	Relation to this work
FEM and adjoint automation [6–10]	High-level variational forms, generated kernels, adjoint-based sensitivities, and PDE-optimization loops.	Provides the reference high-level FEM and optimization context; selected FEniCS runs are scoped comparison formulations.
JAX-native differentiable FEM and PDE solvers [11–14]	Program-level AD, implicit differentiation, Hessian-vector products, nonlinear mechanics, and GPU-oriented differentiable simulation.	Motivates local differentiable modeling; the present study instead couples local JAX derivatives to PETSc sparse MPI solves and compares derivative routes.
FEM–JAX and JAX–PETSc bridge architectures [15–17]	AD-enabled local or variational representations combined with host FEM infrastructure or PETSc sparse solvers.	Architectural context for the JAX+PETSc realization; this study evaluates derivative construction, globalization, and preconditioning on the stated benchmark suite.
Mechanics and topology benchmark lineage [1–5, 18–23]	Strength-reduction plasticity, elastoplastic implementation practice, SIMP topology optimization, filtering, and compact parallel topology software.	Defines the scientific problem classes and lineage; the numerical claims remain limited to the implemented endpoint surrogate and reported comparison observables.
Present toolset	JAX local energies and constitutive laws with PETSc vectors, matrices, nonlinear solvers, Krylov methods, and multigrid.	Across the six-family suite, tests element AD, constitutive AD, colored SFD, nonlinear globalization, sparse assembly, and preconditioner policy in the benchmark subsets where each route is implemented.

Table 1: Selected related computational roles for the software and literature families most relevant to this paper. The entries record documented roles rather than subjective scores.

Finite-element and PDE-optimization automation. The original FEniCS project established the high-level finite-element programming model in which weak forms are expressed symbolically and compiled into efficient kernels [6]. The more recent DOLFINx preprint keeps that mathematical abstraction while moving to a data-oriented design with explicit support for modern parallelism, custom kernels, and direct access to lower-level data structures [7]. In PDE-constrained optimization, `dolfin-adjoint` showed that adjoints of high-level transient finite-element programs can be derived automatically from the symbolic problem description [8], and `pyadjoint` documents the FEniCS/Firedrake adjoint framework used to automate such computations in practice [9]. The `cashocs` software paper extends this PDE-optimization thread to shape optimization, optimal control, topology optimization via level sets, and MPI-enabled computations [10]. These works are the natural reference point for expressive finite-element and adjoint-centered

automation. The present paper uses that context while focusing instead on second-order routes, nonlinear globalization, and sparse distributed solver policy.

Differentiable FEM and JAX–solver bridges. The broad AD background is surveyed by Baydin et al. [24], while JAX provides a program-transformation environment for higher-order derivatives, vectorization, and just-in-time compilation in scientific computing [25]. Within finite elements, JAX-FEM demonstrates inverse design, nonlinear constitutive modeling, and GPU-oriented differentiable computations in a JAX-native 3D FE solver [11]. Xue’s later finite-element differentiable-physics paper is a close methodological reference for second-order information because it derives implicit Hessian-vector products and evaluates Newton–CG use of exact Hessian information on finite-element inverse-problem benchmarks [12]. AutoPDEx broadens the JAX-native PDE picture with nonlinear minimizers, implicit differentiation, and optional PETSc integration [13]; JAX-CPFEM adds differentiable crystal plasticity and GPU execution in a recent mechanics application [14].

Bridge architectures ask a complementary question: how can JAX-style differentiation be combined with established finite-element or sparse-solver environments? The Firedrake–JAX bridge uses tangent-linear and adjoint equations derived from the high-level PDE representation, avoiding differentiation through low-level solver iterations [15]. The FEniCSx external-operator work of Latyshev et al. [16] uses algorithmic automatic differentiation for general constitutive models inside FEniCSx, including a Mohr–Coulomb example. The JetSCI preprint is especially close architecturally because it uses JAX for local finite-element discretization and PETSc for robust distributed sparse solves in heterogeneous micromechanics [17]. These works motivate the primary JAX+PETSc realization studied here. The distinction is that we compare element AD, constitutive AD, and colored sparse finite-difference recovery inside one nonlinear energy-minimization suite spanning the benchmark families below.

Mechanics and topology benchmark context. The plasticity benchmarks use Mohr–Coulomb-type slope-stability problems, so the relevant literature separates continuum plasticity, strength-reduction practice, constitutive integration, and endpoint comparisons. Davis’ plasticity treatment and the finite-element strength-reduction studies of Tschuchnigg et al. [1] provide the Davis-style reduction context used in the paper. For constitutive integration, the incremental reference frame is return mapping with consistent tangents [26, 27], with Mohr–Coulomb-specific nonsmooth return operators treated by Sysala et al. [2]. The recent Sysala slope-stability line documents variational and optimization viewpoints for shear strength reduction [3, 4] and a 3D continuation/iterative-solver reference line [5]. The Čermák–Sysala–Valdman MATLAB implementation [18] documents practical 2D/3D elastoplastic finite-element algorithms for this problem class. We use these works to position the benchmark family; the numerical claims remain limited to the implemented endpoint surrogate and the reported comparison observables.

For topology optimization, the classical SIMP and educational references supply the compliance, filtering, and continuation ideas used throughout the field [19–22]. A relevant modern software reference is FEniTop, which documents a 2D/3D FEniCSx topology-optimization implementation with parallel support [23]. It does not duplicate the reduced design objective used here, but it contextualizes compact parallel topology optimization in the modern FEniCSx ecosystem.

Sparse derivative recovery and scalable nonlinear infrastructure. When exact second-order information is too expensive, sparse derivative recovery via graph coloring remains the classical alternative. Curtis, Powell, and Reid gave the foundational sparse-Jacobian viewpoint [28], while Coleman and Moré extended the approach to sparse Jacobian and Hessian estimation with graph-coloring structure [29, 30]. PETSc supplies the distributed vectors, matrices, Krylov methods, nonlinear solvers, and multigrid infrastructure needed to turn those derivative objects into scalable distributed algorithms [31, 32]. In the present toolset, PETSc ownership layout, Krylov policy, nonlinear globalization, and preconditioner design are part of the numerical method.

Against that background, the present toolset uses JAX for local element and constitutive differentiation, while PETSc owns sparse distributed algebra and nonlinear solver policy for the benchmark suite below.

3 Methodology

We use a common finite-element notation across all benchmark families. The continuous domain is $\Omega \subset \mathbb{R}^d$, and $V_{h,0} = \text{span}\{\varphi_i\}_{i=1}^n$ is the conforming space of free basis functions after homogeneous essential constraints have been applied. Nonhomogeneous essential data are represented by an affine lift $\bar{u}_h(\theta)$, with $\bar{u}_h = 0$ in the homogeneous cases. A free coefficient vector $x \in \mathbb{R}^n$ defines the finite-element state

$$u_h(x; \theta) = \bar{u}_h(\theta) + \sum_{i=1}^n x_i \varphi_i.$$

When no confusion is possible, we identify u_h with its coefficient vector in algebraic formulas. On element e , $x_e = R_e x$ denotes the local coefficient vector selected by the element restriction R_e , while $u_e(x; \theta)$ denotes the corresponding local element state including any contribution from the lift.

The benchmark families use two scalar functional roles. A load-step potential or endpoint surrogate is written as

$$\Pi_h(x; \theta) = \sum_{e=1}^{n_{\text{el}}} \Phi_e(x_e; \theta_e) - f_{\text{ext}}^\top x, \quad (1)$$

where θ collects load, material, lift, continuation, or strength-reduction parameters, θ_e is the element-local restriction of those data, and f_{ext} is the assembled load vector on the free degrees of freedom. A reduced optimization objective with design variable z is written schematically as

$$\mathcal{J}_h(z) = \mathcal{R}_h(z) + \mathcal{S}_h(z, x(z); \theta(z)), \quad (2)$$

where z is a design variable, $x(z)$ is the associated state determined by the benchmark state equation or stationarity condition, and $\mathcal{R}_h$ contains regularization or design penalties. The term $\mathcal{S}_h$ denotes the state-dependent scalar part; in topology optimization it is the frozen compliance term defined in section 5.6, not the load-step potential equation (1). Thus equation (1) covers the scalar PDE, mechanics, and endpoint-surrogate solves, while equation (2) records the reduced-objective role. When an algorithm is stated independently of a benchmark family, we denote its merit functional by $\mathcal{F}$. Hyperelasticity is the only family in which $J = \det F$ denotes the deformation Jacobian, so we avoid the bare symbol J for generic objectives elsewhere in the paper.

In equation (1), Φ_e is the integrated scalar contribution of element e . In quadrature-based mechanics examples, $\Phi_e(x_e; \theta_e) = \sum_q w_q \psi_q(\varepsilon_q; \theta_q)$, where $\varepsilon_q := B_q u_e(x; \theta)$ is the local strain or gradient argument, q indexes quadrature points, w_q are quadrature weights, B_q is the corresponding element operator, θ_q is the quadrature-point restriction of the element data, and ψ_q is the pointwise energy density or branch potential. For finite-element labels, $P_k(L_\ell)$ denotes degree- k Lagrange elements on mesh level L_ℓ ; problem-specific sections define the mesh hierarchy behind those levels when it matters for interpretation.

3.1 Finite-element and derivative structure

The main design choice in the toolset is that problem definitions should remain energy-centered whenever possible. This gives three derivative routes:

1. **element AD:** differentiate the full scalar element contribution $\Phi_e(x_e; \theta_e)$ with respect to the element free-DOF vector;
2. **constitutive AD:** differentiate a quadrature-point energy density $\psi_q(\varepsilon_q, m_q)$ with respect to ε_q , holding branch and material data m_q fixed, then assemble $r_e = \sum_q w_q B_q^\top \sigma_q$ and $K_e = \sum_q w_q B_q^\top C_q B_q$, where $\sigma_q := \partial_\varepsilon \psi_q(\varepsilon_q, m_q)$ and $C_q := \partial_{\varepsilon\varepsilon}^2 \psi_q(\varepsilon_q, m_q)$;
3. **colored sparse recovery:** recover sparse Hessian information through directionally probed gradients or Hessian-vector products (HVPs) and graph coloring.

Element AD is the most literal route and is often the cleanest mathematically. Constitutive AD preserves the same defining scalar-energy expression but changes the differentiation boundary from the full element vector to the local strain state. This is especially effective for high-order plasticity, where the full element Hessian is often much more expensive than the 6×6 constitutive tangent. We use colored sparse recovery as the umbrella term for the colored-SFD configurations reported below: these configurations use AD-HVP probes, while finite-difference gradient probes are the classical variant.

3.2 Newton–Krylov globalization

The nonlinear solves are based on Newton or quasi-Newton updates, but they are not all globalized in the same way. Scalar benchmarks primarily use line-search Newton. Hyperelasticity uses a trust-region method with optional post-step line search. Some benchmark configurations use a hybrid trust-region/line-search algorithm in which a reduced trust-region model provides the direction, while a bounded line search chooses the accepted step. The line-search and trust-region ingredients follow standard globalization ideas [33, Chs. 2–4]. Algorithm 1 is a schematic solver-policy template; the reported runs use the stopping tolerances, caps, globalization choices, and linear-solver/preconditioner choices stated in the benchmark and results tables.

Algorithm 1 Hybrid trust-region / line-search Newton

```
1: Given  $x_0$ , trust radius  $\Delta_0$ , tolerances, and callbacks for energy, gradient, Hessian assembly or application, and linear solves
2: for  $k = 0, 1, \dots$  do
3:   Assemble  $g_k = \nabla \mathcal{F}(x_k)$ 
4:   if convergence criteria are satisfied then
5:     terminate
6:   end if
7:   Approximately solve for a Newton-like direction from  $\nabla^2 \mathcal{F}(x_k)p = -g_k$ 
8:   if trust region is enabled then
9:     Build a small reduced model in the span of Newton and gradient directions
10:    Solve the reduced trust subproblem with radius  $\Delta_k$ 
11:    Optionally compare against a pure gradient trust step
12:    Restrict the line-search interval so the accepted step stays in the trust ball
13:  else
14:    Use the Newton direction directly
15:  end if
16:  Perform the configured one-dimensional line search on the chosen direction
17:  Compute actual and predicted reductions, ared and pred
18:  if trust region is enabled then
19:    accept/reject using  $\rho = \text{ared}/\text{pred}$  and update  $\Delta_k$ 
20:  else
21:    accept according to the configured merit-decrease rule
22:  end if
23: end for
```

Here ared is the achieved decrease in the merit functional and pred is the decrease predicted by the local quadratic or reduced trust-region model. The line-search callback can be Armijo backtracking, a bounded merit search, or a residual-bisection safeguard, depending on the reported configuration. The accept/reject thresholds and radius-update factors are policy parameters; where they affect the interpretation of a result, the paper reports the corresponding configuration rather than treating Algorithm 1 as a complete algorithmic specification.

For the converged bottom-clamped three-dimensional plasticity scaling configuration reported in this paper, the chosen policy is simpler: Newton with Armijo backtracking and a same-mesh polynomial-multigrid $P_4 \rightarrow P_2 \rightarrow P_1$ hierarchy for the flexible-GMRES linear solve [33, Sec. 3.1]. The reported $\lambda_{\text{sr}} = 1.55$ scaling profile uses a redundant LU/MUMPS coarse solve and PETSc relative KSP tolerance 1×10^{-2} ; the auxiliary $\lambda_{\text{sr}} = 1.0$ timing runs use the HyPre-coarse PMG profile. Here λ_{sr} is the plasticity strength-reduction factor defined in section 5.5, PMG denotes same-mesh polynomial multigrid, and KSP denotes PETSc’s Krylov-solver interface and residual tolerance.

Algorithm 2 Armijo line search used in Newton solves

```
1: Given state  $x$ , direction  $p$ , initial step  $\alpha_0$ , backtracking parameters  $\beta \in (0, 1)$ ,  $c_1 \in (0, 1)$ , and maximum count  $m_{\text{max}}$ 
2: Evaluate  $\mathcal{F}(x)$  and  $g = \nabla \mathcal{F}(x)$ 
3: for  $m = 0, 1, \dots, m_{\text{max}}$  do
4:    $\alpha \leftarrow \alpha_0 \beta^m$ 
5:    $x_{\text{trial}} \leftarrow x + \alpha p$ 
6:   if  $\mathcal{F}(x_{\text{trial}}) \leq \mathcal{F}(x) + c_1 \alpha g^\top p$  then
7:     accept  $\alpha$ 
8:     return
9:   end if
10: end for
11: declare line-search failure or use a problem-specific safeguarded steepest-descent fallback when that policy is enabled
```

The numerical tables use a small common reporting vocabulary, summarized in Table 2. These terms are part of the experimental protocol: a completed endpoint solve, a fixed-work diagnostic, and a capped run do not support the same type of claim.

3.3 Colored sparse recovery

Sparse derivative recovery is built around the observation that if columns that do not interfere in the sparsity pattern are perturbed together, one gradient or HVP evaluation can recover multiple columns of the Jacobian or Hessian [28, 30]. The reported colored-SFD path uses distance-2 coloring on the Hessian sparsity graph and AD-HVP probes; the same recovery formula also describes the classical finite-difference gradient-probe variant. This structure is especially effective when combined with overlap-domain assembly and owned-plus-ghost PETSc layouts.

Reported term	Definition in this paper	Interpretation
completed	All configured stopping tests for the configuration are met before the nonlinear cap or wall-time cap.	Endpoint energy and observables are terminal values for the stated benchmark.
timeout / iteration cap	The wall-time limit or maximum nonlinear-iteration count is reached before the stated stopping tests are met.	Retained as solver-behavior evidence, not as validation evidence.
fixed work	A prescribed amount of work is run, such as one Newton step, one linearization, or a fixed continuation schedule.	Used for cost, memory, or policy diagnostics rather than endpoint convergence claims.
wall time	Elapsed end-to-end time for the stated run or run segment, including setup when the reporting path includes it.	Comparable only for configurations that share the same benchmark contract and timing scope.
solve time / solver total	Timer internal to the nonlinear or PETSc solve; setup, state export, and postprocessing may be outside this timer.	Used for within-table solver comparisons, not as a universal replacement for wall time.
relative correction	Norm of the accepted Newton correction relative to the current iterate scale; paired gradient checks use the absolute or relative target stated with the result.	Defines the stopping metric in Plasticity3D rows that report correction or gradient targets.

Table 2: Solver-status and timing vocabulary used in the numerical tables.

Let $H := \nabla^2 \mathcal{F}(x)$ denote the Hessian of the merit functional at the current state, and let $\mathcal{P}(H)$ be its sparsity pattern. A distance-2 coloring partitions the column indices into groups $\mathcal{C}_c$ such that, for each row i , at most one column $j \in \mathcal{C}_c$ satisfies $(i, j) \in \mathcal{P}(H)$. For the color seed $v_c = \sum_{j \in \mathcal{C}_c} e_j$, one probe produces

$$y_c = \begin{cases} (\nabla \mathcal{F}(x + \delta v_c) - \nabla \mathcal{F}(x)) / \delta, & \text{finite-difference gradient probe,} \\ H v_c, & \text{AD Hessian-vector probe.} \end{cases} \quad (3)$$

In the finite-difference branch, $\delta > 0$ is the chosen gradient-probe step size. The structural separation then recovers $H_{ij} = (y_c)_i$ for $j \in \mathcal{C}_c$ and $(i, j) \in \mathcal{P}(H)$. The reported JAX+PETSc colored-SFD cases use the AD Hessian-vector probe mode, so no finite-difference perturbation parameter is introduced in those runs; the finite-difference formula states the classical recovery variant. After recovery, each rank scatters only owned sparse entries to the PETSc matrix, with overlap data used only for local probe evaluation.

Algorithm 3 Distributed colored Hessian recovery

- 1: Build or import local sparsity information for the free-DOF graph
 - 2: Compute a distance-2 coloring of the Hessian pattern
 - 3: **for** each color group c **do**
 - 4: Build the combined probe vector v_c
 - 5: Evaluate one gradient difference or HVP for v_c
 - 6: Scatter recovered entries into owned sparse storage
 - 7: **end for**
 - 8: Assemble the global PETSc matrix from owned COO/CSR entries
-

3.4 Constitutive autodiff for 3D plasticity

The main Plasticity3D result in this paper uses constitutive AD. Rather than applying second-order AD to the full P_4 element vector, we define one scalar Mohr–Coulomb branch potential $\psi_q(\varepsilon_q, m_q)$ per quadrature point, where m_q collects the local material parameters. We differentiate this scalar potential with respect to the six engineering-strain components and assemble the element tangent from the resulting stress and constitutive tangent. This keeps the scalar-energy definition intact while reducing the cost of branchwise second-order information on high-order tetrahedra.

Algorithm 4 Constitutive AD assembly for high-order plasticity

- 1: **for** each local element e **do**
 - 2: **for** each quadrature point q **do**
 - 3: Evaluate strain $\varepsilon_q = B_q u_e(x; \theta)$
 - 4: Evaluate scalar energy density $\psi_q(\varepsilon_q, m_q)$
 - 5: Obtain stress $\sigma_q = \partial_{\varepsilon} \psi_q(\varepsilon_q, m_q)$
 - 6: Obtain tangent $C_q = \partial_{\varepsilon \varepsilon}^2 \psi_q(\varepsilon_q, m_q)$
 - 7: **end for**
 - 8: Assemble $r_e = \sum_q w_q B_q^\top \sigma_q$
 - 9: Assemble $K_e = \sum_q w_q B_q^\top C_q B_q$
 - 10: **end for**
-

The key methodological point is that automatic differentiation is applied directly to the scalar branch potential defined by the benchmark. The constitutive tangent is therefore obtained from that same scalar potential rather than replaced by a separate hand-derived formula. Away from branch switching sets this yields branchwise exact derivatives

of the surrogate; at switching sets the model remains nonsmooth, so the resulting tangents should be read as branchwise surrogate derivatives rather than globally smooth continuum derivatives.

4 Implementation

The computational realization is organized around three solver formulations with different scientific roles.

Pure JAX reference implementations. These are compact differentiable realizations of selected benchmark energies. They are useful for formulation checks, serial references, and AD diagnostics. They are not intended to be the primary scalable route for the largest distributed benchmarks. Their main scientific role in this paper is to show that the same energy formulations can also be expressed without sparse distributed infrastructure.

FEniCS reference implementations. The FEniCS implementations remain valuable on families where they exist, especially p -Laplace, Ginzburg–Landau, and hyperelasticity. They provide a trusted high-level FEM reference and help distinguish modeling issues from solver-policy issues. In this paper they are used mainly as scoped reference comparisons rather than as the primary distributed realization.

Primary JAX+PETSc realization. This is the central solver realization of the paper. It combines JAX-based local physics with PETSc distributed vectors, matrices, Krylov solvers, and preconditioners. Problem-specific kernels stay in JAX; sparse distributed assembly, nonlinear globalization, and linear algebra stay in PETSc.

4.1 Autodiff modes

The JAX+PETSc realization exposes several second-order modes. The precise available choices depend on the benchmark family, but the common logic is:

- **element AD:** local Hessians obtained by differentiating the full element energy on the active smooth branch;
- **constitutive AD:** branchwise local constitutive tangents assembled from quadrature-point energy densities;
- **colored sparse recovery:** colored recovery from AD-HVP or finite-difference gradient probes when full element Hessians are not the appropriate cost tradeoff; tables label this route as colored SFD.

This split is scientifically important. It lets us compare AD-derived local second-order information against carefully structured sparse approximations without changing the nonlinear solver or sparse assembly path. In other words, the paper compares derivative strategies inside one common distributed solver construction, not only across reference implementations.




Figure 1: Computational structure used in the study. The JAX+PETSc formulation is the distributed realization; pure JAX and FEniCS provide reference formulations where matched benchmark definitions are available.




Figure 2: Derivative pathways from scalar energies and quadrature-point potentials to distributed PETSc nonlinear solves.

Element AD $\Pi_e(u_e)$; active-branch local gradient and Hessian; higher-order elements can be costly**Constitutive AD** $\psi_q(\varepsilon_q, m_q)$; quadrature-point tangent assembly; branchwise AD constitutive tangents**Colored SFD**

rank-local AD-HVP probes in reported cases; parallel coloring / recovery; finite-difference gradient probes as the classical variant

Figure 3: Second-order derivative modes exposed in the JAX+PETSc realization. The methodological choice is that the scalar energy or energy-density expression remains the defining mathematical object even when the tangent route changes from element AD to constitutive AD or colored SFD.

Figures 1 to 3 summarize the computational construction used in the rest of the paper. The primary JAX+PETSc realization is the distributed nonlinear solve; the reference formulations provide matched comparisons; and the derivative modes differ only in where the second-order information is constructed before PETSc receives sparse algebraic objects.

4.2 Linear solvers and preconditioners

The benchmark families use different Krylov and preconditioner combinations because the linear systems are qualitatively different. Here CG denotes conjugate gradients, GMRES generalized minimal residual, FGMRES flexible GMRES, GAMG PETSc’s geometric-algebraic multigrid preconditioner, Hypre an algebraic multigrid/preconditioning backend, MUMPS a multifrontal sparse direct solver, and PMG polynomial multigrid:

- **p -Laplace:** mostly CG with Hypre, with exact element Hessians or colored-SFD Hessians in the JAX+PETSc path.
- **Ginzburg–Landau:** GMRES with Hypre, reflecting the more indefinite linearization.
- **Hyperelasticity:** trust-region outer iterations with GAMG-preconditioned Krylov solves in the main benchmark and globalization comparison cases, and same-hierarchy PMG with redundant LU/MUMPS coarse solves in the larger fixed-work scaling profile.
- **Plasticity (2D):** same-mesh multigrid hierarchies whose bottlenecks are dominated by the finest-level smoother and repeated Krylov work.
- **Plasticity3D:** same-mesh polynomial multigrid hierarchies, $P_4 \rightarrow P_2 \rightarrow P_1$ or $P_2 \rightarrow P_1$ depending on degree. The auxiliary $\lambda_{\text{sr}} = 1.0$ scaling and degree/resolution endpoint studies use FGMRES with same-mesh PMG and a CG–Hypre coarse solve, whereas the P_2 globalization and derivative-route comparisons and the converged $P_4(L_2)$, $\lambda_{\text{sr}} = 1.55$ scaling use a redundant LU/MUMPS coarse profile.
- **Topology:** flexible GMRES with GAMG for the mechanics step, coupled to a separate distributed design update.

The result sections also include a fixed-reference PMG variant and alternative Krylov/preconditioner diagnostics. Those are not the primary defaults, but they are informative because they reveal how much of the final performance difference comes from derivative paths, preconditioners, or outer solver logic.

4.3 Distributed assembly and coloring

The JAX+PETSc implementation uses PETSc-owned ranges, local-plus-ghost layouts, and owned-row sparse assembly. Each MPI rank owns a contiguous subset of free degrees of freedom and stores overlap values for elements touching that owned set. Before energy, gradient, Hessian, or HVP work, owned coefficients are scattered into the local ghost layout; the JAX kernels then evaluate element and quadrature contributions rank-locally on this filled overlap state. Matrix entries are inserted only for owned rows, with ghost data used to evaluate local stencils but not to duplicate ownership. Scalar energies are summed from rank-local contributions through PETSc/MPI reductions. This communication order matters for both exact and approximate second-order information. Exact element or constitutive assembly can write directly into local sparse buffers. Colored SFD requires an additional graph-coloring step, but once the coloring is known the probe evaluations can still be carried out in a rank-local manner with only scalar reductions or owned sparse insertion in the dominant assembly loop.




Figure 4: Globalization choices used across the benchmark suite, from pure line search to trust-region and hybrid strategies.




Figure 5: Schematic view of the colored sparse-recovery logic used by the colored-SFD Hessian path: a sparsity pattern, a distance-2 coloring, and the corresponding grouped probes used for recovery.

Figures 4 and 5 make two solver-policy choices explicit. Globalization is varied independently of the element physics where the benchmark permits it, while colored SFD uses graph structure to batch directional Hessian probes without changing the nonlinear problem definition.

Family	FEniCS	pure JAX	JAX+PETSc	Solver and derivative components
p -Laplace	yes	yes	yes	element AD and colored sparse recovery; Newton–CG with Hypre
Ginzburg–Landau	yes	no	yes	element AD and colored sparse recovery; Armijo Newton with GMRES–Hypre
Hyperelasticity	yes	yes	yes	element AD in the primary route, scoped colored-recovery comparison, and trust-region/GAMG or PMG diagnostics
Plasticity2D	no	no	yes	endpoint potential and continuation diagnostics with same-mesh PMG smoother policy
Plasticity3D	no	no	yes	constitutive AD, element AD, and colored-recovery diagnostics; FGMRES with same-mesh PMG and Hypre or LU/MUMPS coarse profiles
Topology optimization	no	yes	yes	distributed design updates with PETSc mechanics and GAMG-preconditioned FGMRES

Table 3: Solver and derivative components by benchmark family. “Solver and derivative components” summarizes the derivative routes, nonlinear policies, or coupled solver features that are scientifically relevant in each family.

Table 3 summarizes the solver components used in the benchmark suite. The most important structural point is that the JAX+PETSc column is the distributed algebraic realization, even when a family also has FEniCS or pure JAX reference formulations. The table is intentionally a component map rather than a full parameter schedule; run-specific tolerances, caps, timing scopes, and coarse-solver variants are reported with the numerical evidence they qualify.

5 Benchmark Problems

The benchmark suite is intentionally heterogeneous. The central claim of the paper is not that one solver policy dominates every nonlinear PDE family, but that one differentiable sparse-solver toolset can cover scalar elliptic problems, phase-field-type energies, finite-strain mechanics, elastoplastic slope stability, and reduced topology optimization without switching conceptual formulations.

Family	Grid / mesh	Solve policy	Compared formulations	Main difficulty
p -Laplace	L_9	Newton + line search	FEniCS, pure JAX, JAX+PETSc	nonlinear elliptic solve with exact sparse Hessians
Ginzburg–Landau	L_9	Newton + line search	FEniCS, JAX+PETSc	indefinite local curvature from the double well
Hyperelasticity	L_4 , 24 steps	trust-region solve	FEniCS, pure JAX, JAX+PETSc	nonconvex large-deformation mechanics
Plasticity2D	$P_4(L_5)$ – $P_4(L_7)$	endpoint solve or fixed nonlinear work	JAX+PETSc only	same-mesh PMG and nonlinear tail behavior
Plasticity3D	$P_1/P_2/P_4$	endpoint, scaling, and diagnostic PMG policies	constitutive and reference PMG variants	heterogeneous 3D Mohr–Coulomb with constitutive AD
Topology	768×384	adaptive continuation	pure JAX, JAX+PETSc	distributed design-mechanics coupling

Table 4: Benchmark families, solve policies, and principal comparison paths. The later numerical section revisits the same families from the timing and scaling viewpoint.

Family	FEniCS	pure JAX	Notes
p -Laplace	yes	yes	All three implementations exist on representative reported cases.
Ginzburg–Landau	yes	no	FEniCS and JAX+PETSc form the reported comparison.
Hyperelasticity	yes	yes	pure JAX is a serial formulation reference only.
Plasticity2D	no	no	The reported realization is JAX+PETSc only.
Plasticity3D	no	no	Reference-formula assembly exists as a supporting comparison route.
Topology	no	yes	Parallel fine-grid realization is JAX+PETSc; pure JAX remains the serial design reference.

Table 5: Availability of FEniCS and pure JAX references. Reference implementations are used where documented implementations exist for the same benchmark family; absent cells are left absent rather than filled with speculative baselines.

Tables 4 and 5 define the comparison scope before the individual models are introduced. The scalar problems and hyperelasticity have FEniCS reference formulations, while pure JAX is used only where a compact serial formulation is available. The two plasticity families are primarily JAX+PETSc studies: the two-dimensional case supplies supporting continuation diagnostics, and the three-dimensional case uses the Čermák–Sysala–Valdman slope-stability implementation lineage as endpoint-surrogate context rather than as a path-history validation baseline. Topology optimization combines a parallel JAX+PETSc realization with a serial pure JAX design reference. Thus some entries below are validation comparisons, whereas others are timing or solver-policy diagnostics under a stated benchmark definition. Table 6 records the compact level notation used in the model definitions and numerical tables; the labels identify discrete problem size and mesh construction, not separate physical models.

Family	Level label	Construction	Representative sizes used below
p -Laplace	L_ℓ	Uniform triangular hierarchy on the L-shaped domain.	L_9 : 788,481 nodes, 1,572,864 triangles, 784,385 free scalar DOFs; L_{10} : 3,149,825 nodes, 6,291,456 triangles, 3,141,633 free scalar DOFs.
Ginzburg–Landau	L_ℓ	Uniform triangular hierarchy on the square domain.	L_9 : 1,050,625 nodes, 2,097,152 triangles, 1,046,529 free scalar DOFs; L_{10} : 4,198,401 nodes, 8,388,608 triangles, 4,190,209 free scalar DOFs.
Hyperelasticity	L_ℓ	Tetrahedral cantilever hierarchy with vector displacement DOFs.	L_4 : 185,249 nodes, 983,040 tetrahedra, 554,013 free displacement DOFs; L_5 : 1,395,009 nodes, 7,864,320 tetrahedra, 4,178,493 free displacement DOFs.
Plasticity2D	$P_4(L_\ell)$	Degree-four same-mesh Lagrange space on the homogeneous slope hierarchy.	$P_4(L_5)$ – $P_4(L_7)$ have 332,960, 1,331,520, and 5,325,440 free displacement DOFs on 20,800, 83,200, and 332,800 triangles.
Plasticity3D	$P_k(L_\ell)$	Degree- k same-mesh Lagrange space on the base heterogeneous tetrahedral slope and one, two, or three uniform refinements.	The P_1 levels L_1 – L_4 have 10,526, 79,024, 610,964, and 4,801,816 free displacement DOFs; P_2 and P_4 use the same macro tetrahedra with elevated element degree.

Table 6: Meaning of the mesh-level labels used in the benchmark and results tables. The counts identify the discrete problem sizes behind the paper-facing labels; they are not additional numerical results.

5.1 p -Laplace

The literature model is the standard stationary p -Laplace Dirichlet problem and its associated energy formulation [34, Ch. 2]. The unknown is a scalar field $u : \Omega \rightarrow \mathbb{R}$ and the continuous energy is

$$\mathcal{E}(u) = \int_{\Omega} \frac{1}{p} |\nabla u|^p dx - \int_{\Omega} f u dx, \quad (4)$$

whose Euler–Lagrange equation is

$$-\nabla \cdot (|\nabla u|^{p-2} \nabla u) = f \quad \text{in } \Omega, \quad u = 0 \quad \text{on } \partial\Omega. \quad (5)$$

The weak problem is: find $u \in W_0^{1,p}(\Omega)$ such that

$$\int_{\Omega} |\nabla u|^{p-2} \nabla u \cdot \nabla v dx = \int_{\Omega} f v dx \quad \forall v \in W_0^{1,p}(\Omega). \quad (6)$$

The benchmark used in this paper takes the L-shaped domain $\Omega = (0, 2)^2 \setminus [1, 2]^2$, exponent $p = 3$, forcing $f = -10$, and a conforming P_1 triangular hierarchy. Here and below, $\mathcal{E}_h$ denotes the finite-element restriction of $\mathcal{E}$ evaluated with the quadrature used for the numerical comparisons. The discrete problem is

$$u_h \in V_h \subset W_0^{1,p}(\Omega) : \quad u_h = \arg \min_{v_h \in V_h} \mathcal{E}_h(v_h), \quad (7)$$

where V_h is the chosen conforming finite-element subspace. This is the cleanest benchmark in the paper for checking exact gradients, exact Hessians, and sparse-recovery alternatives against each other without constitutive or geometric complications.




Figure 6: Representative p -Laplace state on the selected hierarchy.




Figure 7: Final discrete energies from equation (7) across the selected mesh hierarchy. The element-AD and colored-SFD JAX+PETSc paths are energy-consistent with the FEniCS reference at the plotted scale.

Solver realization	Energy	Newton iters	Krylov iters	Wall time [s]
FEniCS Newton reference	-7.942968719	5	15	0.0782
JAX+PETSc element AD	-7.942968719	6	17	0.0798
JAX+PETSc colored SFD	-7.942968711	6	6	0.194

Table 7: Representative final discrete energies from equation (7) on the selected case.

The within-benchmark picture is therefore one of formulation agreement: figure 7 and table 7 show that the implemented routes agree to the reported precision for the discrete energy in equation (7). Scaling and timing for this family are reported in Section 7.3.

5.2 Ginzburg–Landau

We cite Ginzburg and Landau [35] for the historical origin of Ginzburg–Landau free-energy modeling only. The benchmark used here is a scalar real-valued double-well problem on $\Omega = [-1, 1]^2$ with homogeneous Dirichlet boundary conditions and energy

$$\mathcal{E}(u) = \int_{\Omega} \frac{\varepsilon}{2} |\nabla u|^2 + \frac{1}{4} (u^2 - 1)^2 dx, \quad \varepsilon = 10^{-2}. \quad (8)$$

Its strong form is

$$-\varepsilon \Delta u + u(u^2 - 1) = 0 \quad \text{in } \Omega, \quad u = 0 \quad \text{on } \partial\Omega, \quad (9)$$

and the weak problem is: find $u \in H_0^1(\Omega)$ such that

$$\int_{\Omega} \varepsilon \nabla u \cdot \nabla v + (u^2 - 1)uv dx = 0 \quad \forall v \in H_0^1(\Omega). \quad (10)$$

The discrete problem is a conforming P_1 stationary problem on the same triangular hierarchy:

$$u_h \in V_h : \quad D\mathcal{E}_h(u_h)[v_h] = 0 \quad \forall v_h \in V_h. \quad (11)$$

Relative to p -Laplace, the additional difficulty here is not geometry but indefinite local curvature from the double-well potential in equation (8). All compared routes use the same initial state

$$u_0(x, y) = \sin\left(\frac{\pi(x-1)}{2}\right) \sin\left(\frac{\pi(y-1)}{2}\right),$$

restricted to the free degrees of freedom, so the nonconvex branch selection is part of the benchmark definition.




Figure 8: Representative Ginzburg–Landau state on the selected hierarchy.




Figure 9: Final discrete energies associated with the stationary problem equation (11) across mesh levels. The element-AD and colored-SFD JAX+PETSc routes agree with the FEniCS reference to the displayed precision.

Solver realization	Energy	Newton iters	Krylov iters	Wall time [s]
FEniCS Newton reference	0.3462323651	8	33	0.113
JAX+PETSc element AD	0.3462314438	8	33	0.0949
JAX+PETSc colored SFD	0.3462314438	8	33	0.247

Table 8: Representative final discrete energies associated with equation (11) on the selected case.

The local picture is again one of formulation consistency, now for a more solver-sensitive nonconvex energy. Strong scaling and solver-time comparisons for this family are reported in Section 7.4.

5.3 Hyperelasticity

The literature model is a standard finite-strain hyperelastic equilibrium problem written in terms of a displacement u and deformation $y_u(X) = X + u(X)$. We define $F(u) = I + \nabla u$, $J(u) = \det F(u)$, and the compressible Neo-Hookean-type density

$$W(F) := C_1 (\text{tr}(F^\top F) - 3 - 2 \ln \det F) + D_1 (\det F - 1)^2.$$

The first Piola stress is $P(F) = \partial W / \partial F$ [36, Chs. 4–6]. The implemented benchmark used in this paper is a cantilever problem with the potential, defined on admissible displacements with $J(u) > 0$ almost everywhere,

$$\Pi(u) = \int_{\Omega} W(F(u)) \, dx, \quad (12)$$

on $\Omega = [0, 0.4] \times [-0.005, 0.005]^2$ with $C_1 \approx 3.846 \times 10^7$ and $D_1 \approx 8.333 \times 10^7$. The left face is clamped, while the right face carries the prescribed rotating Dirichlet displacement path used throughout this benchmark. Let Γ_L and Γ_R

denote the left and right end faces, respectively. For a reference point $X = (X_1, X_2, X_3)$ on Γ_R and $N = 24$ reported load steps, the prescribed position at step k is

$$\bar{y}_k(X) = (X_1, \cos \eta_k X_2 + \sin \eta_k X_3, -\sin \eta_k X_2 + \cos \eta_k X_3), \quad \eta_k = \frac{8\pi k}{N}.$$

Thus admissible displacements satisfy $u = 0$ on Γ_L and $X + u = \bar{y}_k(X)$ on Γ_R . Let $V(\eta_k)$ denote the corresponding affine displacement space, let V_0 denote the homogeneous test space, and let $\Gamma_N = \partial\Omega \setminus (\Gamma_L \cup \Gamma_R)$. The strong form is

$$-\nabla \cdot P(F(u)) = 0 \quad \text{in } \Omega, \quad u = 0 \quad \text{on } \Gamma_L, \quad X + u = \bar{y}_k(X) \quad \text{on } \Gamma_R, \quad P(F(u))n = 0 \quad \text{on } \Gamma_N, \quad (13)$$

where $P = \partial W / \partial F$ is the first Piola stress. The weak form is: find $u \in V(\eta_k)$ such that

$$\int_{\Omega} P(F(u)) : \nabla v \, dx = 0 \quad \forall v \in V_0. \quad (14)$$

Trial states with nonpositive determinant fall outside the logarithmic energy domain and are rejected by the nonlinear globalization; accepted line-search or trust-region steps are therefore evaluated only on states with positive $J(u)$. The discrete problem solved in the paper is the sequence of stationary points

$$u_h^{(k)} \in V_h(\eta_k) : \quad D_u \Pi_h(u_h^{(k)}; \eta_k)[v_h] = 0 \quad \forall v_h \in V_{h,0}, \quad k = 1, \dots, 24, \quad (15)$$

with first-order tetrahedral vector elements and the benchmark-specific boundary data described above. This is the first family in which globalization is more scientifically important than derivative parity alone.




Figure 10: Representative finite-strain hyperelastic state, colored by the energy density in equation (12).




Figure 11: Final discrete energies along the stationary sequence equation (15) across the selected mesh hierarchy. The comparison is less about exact equality than about tracking the same 24-step trajectory consistently.

Solver realization	Energy	Steps	Krylov iters	Wall time [s]
FEniCS Newton reference	197.775074	24	8722	6.51
JAX+PETSc element AD	197.754582	24	10595	9.36
pure-JAX serial reference	197.749509	24	2284	41.5

Table 9: Representative 24-step terminal energies along equation (15).

The within-benchmark takeaway is therefore that the 24-step stationary sequence is followed consistently: figures 10 and 11 show the terminal deformed state and terminal energies, while globalization and scaling evidence are deferred to Section 7.5. The separate validation section adds one serial JAX-FEM comparison on the same mesh and prescribed displacement schedule, with energy post-comparison under equation (12); it is included as an agreement-focused hyperelastic implementation comparison rather than as a broad software ranking exercise.

5.4 Plasticity2D

The literature model is a quasi-static plane-strain Mohr–Coulomb slope stability problem posed as equilibrium together with a constitutive update and history variables [1–3, 26, 27, 37]. The unknown is the displacement field $u = (u_x, u_y) : \Omega \rightarrow \mathbb{R}^2$ on a homogeneous slope geometry. At increment $n + 1$, the continuum equilibrium equations are

$$-\nabla \cdot \sigma^{n+1} = b^{n+1} \quad \text{in } \Omega, \quad u_x^{n+1} = \bar{u}_x^{n+1} \quad \text{on } \Gamma_{D,x}, \quad u_y^{n+1} = \bar{u}_y^{n+1} \quad \text{on } \Gamma_{D,y}, \quad \sigma^{n+1} n = t^{n+1} \quad \text{on } \Gamma_N, \quad (16)$$

where σ^{n+1} and the internal variables are obtained from the elastoplastic update driven by $\varepsilon(u^{n+1})$ and the previous state. The computational slope has the standard homogeneous benchmark geometry specified by lower bench width 15, slope run 10, upper bench width 15, two vertical height segments of length 10, and slope angle 45° . The bottom edge is fixed in both displacement components, the left and far-right vertical sides are fixed horizontally, and the sloping top is traction-free. The material parameters are $E = 40,000$, $\nu = 0.30$, $c_0 = 6.0$, $\phi = 45^\circ$, $\psi = 0^\circ$, and gravity magnitude $\gamma = 20$, so $b = (0, -\gamma)^\top$ in the endpoint solves. The endpoint studies set the old plastic strain to zero at all quadrature points. Let Γ_{bot} denote the bottom edge and let Γ_{side} denote the constrained vertical side edges. We use $\Gamma_{D,x} = \Gamma_{\text{bot}} \cup \Gamma_{\text{side}}$ and $\Gamma_{D,y} = \Gamma_{\text{bot}}$, so the homogeneous test space V_0 contains functions with $v_x = 0$ on $\Gamma_{D,x}$ and $v_y = 0$ on $\Gamma_{D,y}$. The traction boundary is $\Gamma_N = \partial\Omega \setminus (\Gamma_{\text{bot}} \cup \Gamma_{\text{side}})$. The corresponding weak form is

$$\int_{\Omega} \sigma^{n+1} : \varepsilon(v) dx = \int_{\Omega} b^{n+1} \cdot v dx + \int_{\Gamma_N} t^{n+1} \cdot v ds \quad \forall v \in V_0. \quad (17)$$

The runs reported here are endpoint solves, not path-consistent incremental histories. Accordingly, the benchmark uses the following scalarized discrete surrogate functional for differentiation and comparison:

$$\Pi_h^{2D}(u_h; \lambda_{\text{sr}}) = \sum_{e=1}^{n_{\text{el}}} \sum_{q=1}^{n_q} w_{eq} \Phi_{\text{MC},2D}(B_{eq} u_e, \varepsilon_{eq}^{p,\text{old}}, E, \nu, c_{\lambda,eq}, \phi_{\lambda,eq}) - f_{\text{ext}}^\top u_h, \quad (18)$$

where λ_{sr} denotes the strength-reduction factor, distinct from the Lamé parameter used later in three dimensions, and $\Phi_{\text{MC},2D}$ denotes the scalar branch potential used for the two-dimensional Mohr–Coulomb endpoint surrogate. Here the subscript eq means element e and quadrature point q , w_{eq} is the corresponding quadrature weight including the element Jacobian, and B_{eq} maps the element displacement vector to the two-dimensional engineering-strain vector used by the branch potential. The load vector f_{ext} is assembled from the right-hand side of equation (17); for the endpoint solves this is the gravity load with $b = (0, -\gamma)^\top$ and zero traction on the free boundary. The dilation angle enters this endpoint model only through the Davis-B reduction below. The reduced cohesion and friction angle appearing in equation (18) are defined by the Davis-B reduction equation (19).

The two-dimensional branch formulas use engineering strain $\varepsilon = (\varepsilon_{xx}, \varepsilon_{yy}, \gamma_{xy})^\top$ and the plane-strain elastic matrix

$$C_{2D} = \begin{bmatrix} \lambda_{\text{ps}} + 2\mu & \lambda_{\text{ps}} & 0 \\ \lambda_{\text{ps}} & \lambda_{\text{ps}} + 2\mu & 0 \\ 0 & 0 & \mu \end{bmatrix}, \quad \lambda_{\text{ps}} = \frac{E\nu}{(1+\nu)(1-2\nu)}, \quad \mu = \frac{E}{2(1+\nu)}.$$

Its compliance is $S_{2D} = C_{2D}^{-1}$. Principal stresses are ordered as $\sigma_1^{\text{tr}} \geq \sigma_2^{\text{tr}}$, and the elastic branch is the branch for which the trial function in the ordered variables satisfies $f^{\text{tr}} \leq 0$.

The benchmark uses a Davis-B-style reduction path [1, 3, 37] to map the unreduced cohesion and angles to reduced values under strength reduction; with c_0, ϕ, ψ denoting the unreduced cohesion, friction angle, and dilation angle, respectively, the benchmark uses

$$\begin{aligned} c_1 &= c_0 / \lambda_{\text{sr}}, & \phi_1 &= \arctan(\tan \phi / \lambda_{\text{sr}}), & \psi_1 &= \arctan(\tan \psi / \lambda_{\text{sr}}), \\ \beta &= \frac{\cos \phi_1 \cos \psi_1}{1 - \sin \phi_1 \sin \psi_1}, & c_\lambda &= \beta c_1, \\ \phi_\lambda &= \arctan(\beta \tan \phi_1), & \psi_\lambda &= \arctan(\beta \tan \psi_1). \end{aligned} \quad (19)$$

The reported endpoint runs in this paper use $\lambda_{\text{sr}} = 1$ and $\psi = 0^\circ$, hence $\psi_\lambda = 0^\circ$; the reported states are algorithmic endpoint surrogates rather than path-consistent history updates. For a trial engineering strain $\varepsilon = (\varepsilon_{xx}, \varepsilon_{yy}, \gamma_{xy})^\top$ and old plastic strain $\varepsilon^{p,\text{old}}$, set

$$\varepsilon^e = \varepsilon - \varepsilon^{p,\text{old}}, \quad \sigma^{\text{tr}} = C_{2D} \varepsilon^e, \quad S_{2D} = C_{2D}^{-1}.$$

Writing $\sigma^{\text{tr}} = (\sigma_{xx}, \sigma_{yy}, \tau_{xy})^\top$, the principal stresses use

$$\sigma_{1,2}^{\text{tr}} = m \pm \sqrt{d^2 + \tau_{xy}^2 + \epsilon_{\text{reg}}^2}, \quad m = \frac{\sigma_{xx} + \sigma_{yy}}{2}, \quad d = \frac{\sigma_{xx} - \sigma_{yy}}{2}.$$

With $a_\lambda = 1 + \sin \phi_\lambda$, $b_\lambda = 1 - \sin \phi_\lambda$, and $\kappa_\lambda = 2c_\lambda \cos \phi_\lambda$, the trial state is elastic when

$$f^{\text{tr}} = a_\lambda \sigma_1^{\text{tr}} - b_\lambda \sigma_2^{\text{tr}} - \kappa_\lambda \leq 0.$$

In that case $\sigma^* = \sigma^{\text{tr}}$. Otherwise let

$$z = \hat{S}^{-1}(a_\lambda, -b_\lambda)^\top, \quad \eta = \frac{f^{\text{tr}}}{a_\lambda z_1 - b_\lambda z_2}, \quad \hat{\sigma}^\ell = (\sigma_1^{\text{tr}}, \sigma_2^{\text{tr}})^\top - \eta z,$$

where $\hat{S}$ is the principal-stress 2×2 normal block of $S_{2\text{D}}$. If $\hat{\sigma}_1^\ell \geq \hat{\sigma}_2^\ell$, set $\hat{\sigma}^* = \hat{\sigma}^\ell$; otherwise set $\hat{\sigma}_1^* = \hat{\sigma}_2^* = \kappa_\lambda / (a_\lambda - b_\lambda)$. After rotating $\hat{\sigma}^*$ back to Cartesian components, the branch potential is

$$\Phi_{\text{MC},2\text{D}} = \sigma^* \cdot \varepsilon^e - \frac{1}{2} \sigma^* \cdot S_{2\text{D}} \sigma^*.$$

The reported two-dimensional rows use this quantity as a solver-diagnostic endpoint surrogate. The small square-root regularization parameter $\epsilon_{\text{reg}} = 1 \times 10^{-12}$ smooths the endpoint surrogate; it is not a separate material parameter.




Figure 12: Representative Plasticity2D state: displacement magnitude and deviatoric strain derived from the scalar surrogate equation (18).




Figure 13: Resolution trend for the Plasticity2D family. The plotted terminal energies are evaluations of equation (18) on the completed $P_4(L_5)$ endpoint case and on the larger fixed-iteration $P_4(L_6)$ and $P_4(L_7)$ diagnostics, which are not used as endpoint-convergence evidence.

Case	Free DOFs	Energy	Wall time [s]	Note
$P_4(L_5)$	332,960	-212.538	147	completed endpoint
$P_4(L_6)$	1,331,520	-212.539	147	fixed-iteration diagnostic at 8 ranks
$P_4(L_7)$	5,325,440	-212.539	253	fixed-iteration diagnostic at 16 ranks

Table 10: Representative Plasticity2D endpoint and fixed-work diagnostic values from equation (18).

This family is where multigrid hierarchy design first becomes dominant in the paper. Only the completed $P_4(L_5)$ endpoint case is treated as an endpoint solve; the larger $P_4(L_6)$ and $P_4(L_7)$ cases are fixed-iteration diagnostics for solver behavior. Solver-engineering evidence for Plasticity2D is reported later in Sections A and 7.6.

5.5 Plasticity3D

The literature model is again quasi-static Mohr–Coulomb equilibrium with constitutive/history updates, now on a three-dimensional heterogeneous slope geometry [1–5, 26, 27, 37]. The unknown is a displacement field $u = (u_x, u_y, u_z) : \Omega \rightarrow \mathbb{R}^3$, and gravity acts in the vertical y -direction. Let $e_y = (0, 1, 0)^\top$. Let $\Gamma_{D,i}$ denote the component-wise Dirichlet boundary for displacement component i , let Γ_N be the remaining traction-free boundary, and let V_0 be the homogeneous vector-valued test space after applying these component-wise constraints. With material-region weight $\gamma(x)$, the continuum equilibrium equations are

$$-\nabla \cdot \sigma^{n+1} = -\gamma(x)e_y \quad \text{in } \Omega, \quad u_i^{n+1} = \bar{u}_i^{n+1} \quad \text{on } \Gamma_{D,i}, \quad \sigma^{n+1}n = 0 \quad \text{on } \Gamma_N. \quad (20)$$

where σ^{n+1} and the internal variables are obtained from the elastoplastic update driven by $\varepsilon(u^{n+1})$ and the previous state. The heterogeneous benchmark uses four material regions. Their $(c_0, \phi, \psi, E, \nu, \gamma)$ values, in the order cover layer, general foundation, weak foundation, and slope mass, are $(15, 30^\circ, 0^\circ, 10,000, 0.33, 19)$, $(15, 38^\circ, 0^\circ, 50,000, 0.30, 22)$, $(10, 35^\circ, 0^\circ, 50,000, 0.30, 21)$, and $(18, 32^\circ, 0^\circ, 20,000, 0.33, 20)$. At quadrature point q , the body force is $-\gamma_q e_y$, where γ_q is selected by the material region. The elastic constants used below are

$$\mu = \frac{E}{2(1+\nu)}, \quad K = \frac{E}{3(1-2\nu)}, \quad \lambda_L = K - \frac{2\mu}{3}.$$

The corresponding weak form is

$$\int_{\Omega} \sigma^{n+1} : \varepsilon(v) dx = - \int_{\Omega} \gamma(x)e_y \cdot v dx \quad \forall v \in V_0. \quad (21)$$

The studied discrete model uses a heterogeneous tetrahedral mesh with four material physical groups, refined into the displayed mesh hierarchy and elevated locally to same-mesh $P_1/P_2/P_4$ spaces. Let $\Gamma_1, \dots, \Gamma_5$ denote the boundary label sets of the reference mesh. In coordinates $X = (X_1, X_2, X_3)$, the mesh bounding box is $[-175, 30] \times [0, 60] \times [0, 86.6025]$ and X_2 is vertical. The labels $\Gamma_1, \Gamma_2, \Gamma_3, \Gamma_4, \Gamma_5$ correspond, respectively, to the right face $X_1 = 30$, the left face $X_1 = -175$, the front face $X_3 = 0$, the back face $X_3 = 86.6025$, and the bottom face $X_2 = 0$; the remaining exterior labels contain the sloped free surface and the top face. We impose

$$\Gamma_{D,x} = \Gamma_1 \cup \Gamma_2, \quad \Gamma_{D,y} = \Gamma_5, \quad \Gamma_{D,z} = \Gamma_3 \cup \Gamma_4.$$

The bottom-clamped model additionally fixes all displacement components on $\Gamma_{\text{bot}} = \{X \in \partial\Omega : |X_2| \leq 10^{-12}\}$. We write $P_k(L_\ell)$ for polynomial degree k on mesh level ℓ and use $L_1, \dots, L_4$ for the four displayed mesh levels. The figures below use the bottom-clamped $P_4(L_2)$ representative endpoint at $\lambda_{\text{sr}} = 1.55$ and the finished nine-case study at the same load factor. The $\lambda_{\text{sr}} = 1.0$ strong-scaling study is kept later in the paper as separate auxiliary timing evidence. For those endpoint and continuation studies, we evaluate the following algorithmic scalar surrogate:

$$\Pi_h^{3D}(u_h; \lambda_{\text{sr}}) = \sum_{e=1}^{n_{\text{el}}} \sum_{q=1}^{n_q} w_{eq} \Phi_{\text{MC},3D}(\varepsilon_{eq}(u_e), \bar{c}_{\lambda,eq}, \sin \phi_{\lambda,eq}, \mu_{eq}, K_{eq}, \lambda_{L,eq}) - f_{\text{ext}}^\top u_h, \quad (22)$$

where λ_{sr} is again the strength-reduction factor and λ_L is the Lamé parameter of the elastic predictor. The subscript eq means that each material parameter is evaluated at element e and quadrature point q . The load vector f_{ext} is assembled from the body force $-\gamma_q e_y$ in equation (21), with no Neumann traction contribution in the reported cases.

We use the engineering-strain ordering

$$\varepsilon = [\varepsilon_{xx} \quad \varepsilon_{yy} \quad \varepsilon_{zz} \quad \gamma_{xy} \quad \gamma_{yz} \quad \gamma_{xz}]^\top, \quad (23)$$

which is first mapped to tensor shear before principal values are formed. The benchmark keeps the Davis-B reduction explicitly at quadrature points:

$$\begin{aligned} c_{1,q} &= c_{0,q}/\lambda_{\text{sr}}, & \phi_{1,q} &= \arctan(\tan \phi_q/\lambda_{\text{sr}}), & \psi_{1,q} &= \arctan(\tan \psi_q/\lambda_{\text{sr}}), \\ \beta_q &= \frac{\cos \phi_{1,q} \cos \psi_{1,q}}{1 - \sin \phi_{1,q} \sin \psi_{1,q}}, & c_{\lambda,q} &= \beta_q c_{1,q}, & \phi_{\lambda,q} &= \arctan(\beta_q \tan \phi_{1,q}), \end{aligned} \quad (24)$$

with

$$\bar{c}_{\lambda,q} = 2c_{\lambda,q} \cos \phi_{\lambda,q}, \quad \sin \phi_{\lambda,q} = \sin(\phi_{\lambda,q}), \quad (25)$$

which are the scalar constitutive parameters used by the 3D assembly path [1, 3, 37]. Here ϕ denotes friction angle and ψ denotes dilation angle; the paper uses those symbols only in that geomechanics sense. All 3D material tuples used in the reported runs have $\psi = 0^\circ$; the dilation angle therefore enters only through the Davis-B reduction above and not as a separate argument of $\Phi_{\text{MC},3D}$.

The local constitutive surrogate follows the piecewise branch structure of the reference model. At one quadrature point we suppress the index q and write $s_\lambda := \sin \phi_\lambda$. Let $\varepsilon_1 \geq \varepsilon_2 \geq \varepsilon_3$ be the ordered principal strains obtained after

converting engineering shear components to tensor shear. The branch tests below use the convention that the elastic case satisfies $f_{\text{tr}} \leq 0$. Let

$$I_1 = \varepsilon_{xx} + \varepsilon_{yy} + \varepsilon_{zz}, \quad Q(\varepsilon) = \varepsilon_{xx}^2 + \varepsilon_{yy}^2 + \varepsilon_{zz}^2 + \frac{1}{2}(\gamma_{xy}^2 + \gamma_{yz}^2 + \gamma_{xz}^2). \quad (26)$$

The elastic branch is

$$\Phi_{\text{el}} = \frac{1}{2}\lambda_L I_1^2 + \mu Q(\varepsilon), \quad (27)$$

and the branch-selection variables are

$$\begin{aligned} f_{\text{tr}} &= 2\mu((1+s_\lambda)\varepsilon_1 - (1-s_\lambda)\varepsilon_3) + 2\lambda_L s_\lambda I_1 - \bar{c}_\lambda, \\ \gamma_{\text{sl}} &= \frac{\varepsilon_1 - \varepsilon_2}{1+s_\lambda}, & \gamma_{\text{sr}} &= \frac{\varepsilon_2 - \varepsilon_3}{1-s_\lambda}, \\ \gamma_{\text{la}} &= \frac{\varepsilon_1 + \varepsilon_2 - 2\varepsilon_3}{3-s_\lambda}, & \gamma_{\text{ra}} &= \frac{2\varepsilon_1 - \varepsilon_2 - \varepsilon_3}{3+s_\lambda}. \end{aligned} \quad (28)$$

The return denominators are

$$\begin{aligned} D_s &= 4\lambda_L s_\lambda^2 + 4\mu(1+s_\lambda^2), & D_l &= 4\lambda_L s_\lambda^2 + \mu(1+s_\lambda)^2 + 2\mu(1-s_\lambda)^2, \\ D_r &= 4\lambda_L s_\lambda^2 + 2\mu(1+s_\lambda)^2 + \mu(1-s_\lambda)^2, & D_a &= 4K s_\lambda^2. \end{aligned} \quad (29)$$

The corresponding branch multipliers are

$$\begin{aligned} \Lambda_s &= \frac{f_{\text{tr}}}{D_s}, \\ \Lambda_l &= \frac{\mu((1+s_\lambda)(\varepsilon_1 + \varepsilon_2) - 2(1-s_\lambda)\varepsilon_3) + 2\lambda_L s_\lambda I_1 - \bar{c}_\lambda}{D_l}, \\ \Lambda_r &= \frac{\mu(2(1+s_\lambda)\varepsilon_1 - (1-s_\lambda)(\varepsilon_2 + \varepsilon_3)) + 2\lambda_L s_\lambda I_1 - \bar{c}_\lambda}{D_r}, \\ \Lambda_a &= \frac{2K s_\lambda I_1 - \bar{c}_\lambda}{D_a}. \end{aligned} \quad (30)$$

The branch energies are

$$\begin{aligned} \Phi_s &= \Phi_{\text{el}} - \frac{1}{2}D_s \Lambda_s^2, \\ \Phi_l &= \frac{1}{2}\lambda_L I_1^2 + \mu\left(\varepsilon_3^2 + \frac{1}{2}(\varepsilon_1 + \varepsilon_2)^2\right) - \frac{1}{2}D_l \Lambda_l^2, \\ \Phi_r &= \frac{1}{2}\lambda_L I_1^2 + \mu\left(\varepsilon_1^2 + \frac{1}{2}(\varepsilon_2 + \varepsilon_3)^2\right) - \frac{1}{2}D_r \Lambda_r^2, \\ \Phi_a &= \frac{1}{2}K I_1^2 - \frac{1}{2}D_a \Lambda_a^2. \end{aligned} \quad (31)$$

With these definitions, the scalar potential $\Phi_{\text{MC},3\text{D}}(\varepsilon; \bar{c}_\lambda, s_\lambda, \mu, K, \lambda_L)$ is the piecewise function

$$\Phi_{\text{MC},3\text{D}}(\varepsilon; \bar{c}_\lambda, s_\lambda, \mu, K, \lambda_L) = \begin{cases} \Phi_{\text{el}}, & f_{\text{tr}} \leq 0, \\ \Phi_s, & \Lambda_s \leq \min(\gamma_{\text{sl}}, \gamma_{\text{sr}}), \\ \Phi_l, & \gamma_{\text{sl}} < \gamma_{\text{sr}} \text{ and } \gamma_{\text{sl}} \leq \Lambda_l \leq \gamma_{\text{la}}, \\ \Phi_r, & \gamma_{\text{sr}} < \gamma_{\text{sl}} \text{ and } \gamma_{\text{sr}} \leq \Lambda_r \leq \gamma_{\text{ra}}, \\ \Phi_a, & \text{otherwise,} \end{cases} \quad (32)$$

The constitutive-AD path differentiates the smooth branches of equation (32) directly with JAX rather than substituting a hand-derived return-map tangent; the branch interfaces remain nonsmooth and are treated as part of the algorithmic constitutive surrogate rather than as a globally smooth continuum potential.

For the strain visualizations in this paper we use the deviatoric norm

$$\|\varepsilon_{\text{dev}}\| = \sqrt{\varepsilon^\top \left(\text{diag}(1, 1, 1, \frac{1}{2}, \frac{1}{2}, \frac{1}{2}) - \frac{1}{3}mm^\top \right) \varepsilon}, \quad m = (1, 1, 1, 0, 0, 0)^\top, \quad (33)$$

which matches the engineering-shear convention of equation (23).

Because equations (22) and (32) define an algorithmic endpoint surrogate rather than a path-consistent incremental history functional, the results section reports this family with a separate endpoint-surrogate comparison. Direct-branch endpoint-observable differences against the reference model are treated separately from a fixed- λ_{sr} reference comparison.

The latter checks matched geometry, materials, Davis-B reduction parameters, load schedule, and the active boundary conditions, but it is not a path-consistent incremental-history validation. The paper therefore keeps the Plasticity3D claims at the level of a problem-specific endpoint surrogate rather than a path-consistent replacement. The related return-mapping, optimization, convex-optimization, and continuation literature [2–5] is used as reference-model context, not as evidence that this endpoint surrogate is a full incremental elastoplastic solver.




Figure 14: Bottom-clamped Plasticity3D representative endpoint at $\lambda_{\text{sr}} = 1.55$: displacement magnitude and surface deviatoric strain for the $P_4(L_2)$ run, both derived from the scalar surrogate in equation (22).




Figure 15: Newton-history convergence for all nine completed bottom-clamped Plasticity3D study runs at $\lambda_{\text{sr}} = 1.55$. The upper panel shows the value of the algorithmic surrogate from equation (22); the lower panel shows the gradient norm used as the nonlinear stopping metric, with color indicating degree and line style indicating mesh level.

Element	Free DOFs	Energy	$\ g\ _{\text{final}}$	Wall time [s]
$P_1(L_1)$	10,526	-2,953,167	4.56×10^{-3}	1.79
$P_1(L_2)$	79,024	-2,991,497	1.64×10^{-3}	8.72
$P_1(L_3)$	610,964	-3,004,742	4.07×10^{-3}	17.2
$P_1(L_4)$	4,801,816	-3,009,658	1.16×10^{-3}	96.7
$P_2(L_1)$	79,024	-3,012,813	2.47×10^{-3}	18.3
$P_2(L_2)$	610,964	-3,013,032	2.96×10^{-3}	98.3
$P_2(L_3)$	4,801,816	-3,013,149	1.14×10^{-3}	1330
$P_4(L_1)$	610,964	-3,013,227	5.53×10^{-3}	208
$P_4(L_2)$	4,801,816	-3,013,348	4.42×10^{-3}	2298

Table 11: All nine completed bottom-clamped Plasticity3D study cases at $\lambda_{\text{sr}} = 1.55$, reported as endpoint values of equation (22).




Figure 16: Shared-scale upper y -slice comparison of $\|\varepsilon_{\text{dev}}\|$ from equation (33) at the highest mesh level of each degree line: $P_1(L_4)$, $P_2(L_3)$, and $P_4(L_2)$. The three runs use the bottom-clamped solver configuration at $\lambda_{\text{sr}} = 1.55$ and are compared on one common color scale.

Figures 14 to 16 and table 11 establish the benchmark object used later in the results section. The nine completed endpoint solves converge in the nonlinear stopping metric while the terminal values of equation (22) decrease along both mesh and degree refinements. The shared-scale slice comparison shows that the larger values of $\|\varepsilon_{\text{dev}}\|$ remain localized near the same upper slope region across P_1 , P_2 , and P_4 , so the later degree and timing comparisons are interpreted as endpoint-solver comparisons rather than changes of the physical problem.

Plasticity3D is the benchmark where detailed constitutive definitions are most important: the paper later plots values of equation (22) and equation (33), so both quantities are defined explicitly here. Timing, scaling, and degree-versus-resolution evidence for this family are reported in Section 7.7.

5.6 Topology Optimization

The literature model is the standard SIMP-style compliance topology-optimization setting on a cantilever domain, together with filtering and continuation ideas from the classical topology literature [19–22]. In the implemented setting studied here, the topology benchmark is a reduced optimization problem rather than a single equilibrium minimization. Let Z_h be the scalar continuous P_1 space on the structured triangular mesh. The design variable is a latent nodal field $z_h \in Z_h$, mapped to a physical density through the logistic transform

$$\theta_h(z_h) = \theta_{\min} + (1 - \theta_{\min})\sigma(z_h), \quad \sigma(z) = \frac{1}{1 + e^{-z}}. \quad (34)$$

On each triangular element e , we use the element averages $\theta_e := |e|^{-1} \int_e \theta_h dx$ and $z_e := |e|^{-1} \int_e z_h dx$; for the linear elements used here these are the vertex averages. The reported fine-grid run uses a 768×384 mesh, target volume fraction 0.4, $\theta_{\min} = 1 \times 10^{-6}$, Young’s modulus $E = 1$, Poisson ratio 0.3, and a downward unit traction over the centered 20% patch of the right edge. The left edge is clamped, while fixed solid design pads are imposed near the clamp and around the load patch. For the domain $\Omega = (0, 2) \times (0, 1)$ this load patch is $\{2\} \times [0.4, 0.6]$. With $h_x = 2/768$ and $h_y = 1/384$, the fixed-solid design pads are $x \leq 32h_x$ and $x \geq 2 - 32h_x$, $y \in [0.4 - 32h_y, 0.6 + 32h_y]$. Let V_h be the

continuous vector P_1 displacement space and let $V_{h,0} \subset V_h$ impose the homogeneous left-edge clamp. Let C denote the plane-stress elasticity tensor and let p be the active SIMP exponent. For fixed density, define the bilinear form

$$a_{\theta_h}(w_h, v_h) := \sum_{e=1}^{n_{el}} |e| \theta_e^p (\varepsilon(w_h)|_e)^\top C (\varepsilon(v_h)|_e). \quad (35)$$

Here $\varepsilon(v) = (\partial_x v_x, \partial_y v_y, \partial_y v_x + \partial_x v_y)^\top$, and

$$C = \frac{E}{1-\nu^2} \begin{bmatrix} 1 & \nu & 0 \\ \nu & 1 & 0 \\ 0 & 0 & (1-\nu)/2 \end{bmatrix}.$$

The load functional is

$$\ell(v) = \int_{\{2\} \times [0.4, 0.6]} (0, -1) \cdot v \, ds.$$

The mechanics step finds $u_h(\theta_h)$ in the affine displacement space such that $a_{\theta_h}(u_h, v_h) = \ell(v_h)$ for all $v_h \in V_{h,0}$. At outer design iteration m , after solving for $u_h^m = u_h(\theta_h^m)$, the design step freezes, on each element,

$$e_{e,m}^{\text{frozen}} = (\theta_e^m)^{2p} \varepsilon(u_h^m)^\top C \varepsilon(u_h^m).$$

Let $(\nabla \theta_h)|_e$ denote the elementwise gradient of the interpolated density field used in the phase-field term. The reduced objective for the next design vector z_h^{m+1} is

$$\mathcal{J}_h^m(z_h) = \sum_{e=1}^{n_{el}} |e| \left[e_{e,m}^{\text{frozen}} \theta_e^{-p} + \lambda_V \theta_e + \alpha_{\text{reg}} \left(\frac{\ell_{\text{pf}}}{2} |(\nabla \theta_h)|_e|^2 + \frac{1}{\ell_{\text{pf}}} W(\theta_e) \right) + \frac{\mu_{\text{move}}}{2} (z_e - z_e^m)^2 \right], \quad (36)$$

where θ_e and z_e are computed from the trial z_h , while z_e^m is the element average from the previous design iterate. The function $W(\theta) = \theta^2(1-\theta)^2$ supplies the double-well term, while the SIMP-style continuation in p and the regularization terms play a mesh-control role analogous to filtering/regularization in Bourdin [22]. Here λ_V is the effective volume multiplier, updated with penalty weight 10 and multiplier step factor 12, while $\alpha_{\text{reg}} = 5 \times 10^{-3}$, $\ell_{\text{pf}} = 8 \times 10^{-2}$, and $\mu_{\text{move}} = 1 \times 10^{-2}$ set the phase-field and latent-move terms. These are algorithmic design choices rather than standard SIMP notation. The outer loop alternates mechanics solves for equation (35) with reduced design updates for equation (36).




Figure 17: Representative final density field for the parallel topology benchmark.




Figure 18: Objective-history view for the reduced topology objective equation (36), with compliance on the left axis and volume fraction plus the continuation schedule in p on the right axis.

Case	Ranks	Outer iters	Compliance	Volume fraction	Wall time [s]
192×96	1	121	4.1557	0.4	10
768×384	32	66	9.9074	0.3749	25.1

Table 12: Representative topology endpoint values from equation (36).

For topology, the relevant notion of “convergence” is the reduced objective history rather than a single equilibrium energy. The reported continuation starts at SIMP exponent $p = 1$, increases p by 0.2 on each outer iteration up to $p = 10$, and treats the fine-grid run as stalled once $p \geq 4$ and both the design change and state change are at most 10^{-6} . Parallel timing and scaling for the same family are reported in Section 7.8.

6 Validation and External Comparisons

This section separates validation and external-comparison evidence from the performance results that follow. The distinction is important because the paper uses different levels of evidence for different claims: Hyperelasticity has a narrow external comparison against JAX-FEM, while Plasticity3D is evaluated through a reference-model endpoint-surrogate comparison.

6.1 Comparison Metrics

For vector and field comparisons with nonzero reference norm, we report relative differences of the form

$$\delta_M(a, b) = \frac{\|a - b\|_M}{\|a\|_M}, \quad (37)$$


where a is the reference observable, b is the compared observable, and M denotes the discrete norm used for the comparison. For nodal displacement vectors, M is the Euclidean norm on the compared free degrees of freedom. For sampled profiles and quadrature-derived fields, M denotes the corresponding sampled or quadrature-weighted norm used by the comparison. The reported field and profile observables have nonzero reference norms. Scalar work and energy differences use the same reference-denominator convention; the hyperelastic energy post-comparison uses an absolute denominator floor of 10^{-12} .

The hyperelastic JAX-FEM comparison uses a 5% predefined threshold for the final energy, full-field displacement, centerline displacement, and $u_{\max}$ -curve differences. The Plasticity3D fixed-load reference comparison uses three predefined criteria: at most 0.03 relative difference in the highest successful load factor, at most 0.05 relative L^2 difference for the $u_{\max}(\lambda_{\text{sr}})$ curve, and at most 0.10 relative L^2 difference for the endpoint displacement. The deviatoric-strain and boundary-profile differences in Table 14 are diagnostics, not acceptance criteria.

6.2 Hyperelasticity External Comparison

The hyperelasticity energy is defined in Section 5.3. The external comparison here is deliberately narrow and uses a serial companion case rather than the 24-step rotating benchmark path. Both implementations use the same level-1 mesh and the same right-face x -translation schedule $d = (0.0025, 0.005, 0.0075, 0.01)$, then post-compare both terminal

states with the energy from equation (12). This supports a solution-level comparison for this hyperelastic problem only. The comparison is neither a plasticity validation nor a broad software ranking, and it does not assert constitutive-law identity between the two implementations: the JAX-FEM implementation uses a logarithmic- J Piola form, whereas equation (12) uses the $D_1(J - 1)^2$ volumetric penalty.



(a) Terminal energy over the matched displacement schedule.

(b) Final centerline u_x profile on the common sample path.

Figure 19: Serial hyperelastic comparison against JAX-FEM on the same mesh and prescribed displacement schedule, with energy post-comparison of the terminal states under equation (12).

Group	Quantity	Relative difference	Status
Agreement	final energy	8.24×10^{-6}	below 5%
Agreement	full-field displacement relative L^2	2.01×10^{-4}	below 5%
Agreement	centerline relative L^2	5.47×10^{-5}	below 5%
Agreement	$u_{\max}$ curve relative L^2	8.85×10^{-16}	below 5%
Condition	common mesh	—	satisfied
Condition	common displacement schedule	—	satisfied
Condition	agreement threshold (5%)	—	satisfied
Condition	energy re-evaluation	—	applied

Table 13: Thresholded hyperelastic agreement comparison against JAX-FEM.


The agreement criteria in Table 13 are satisfied for all checked quantities. The final energy relative difference is 8.244×10^{-6} , the full-field displacement relative L^2 difference is 2.011×10^{-4} , the centerline relative L^2 difference is 5.473×10^{-5} , and the $u_{\max}$ -curve relative L^2 difference is effectively zero. The scientific role of this table is terminal-state agreement under a matched mesh, displacement schedule, and common energy post-comparison rather than a software ranking.

6.3 Plasticity3D Endpoint-Surrogate Comparison


The Plasticity3D benchmark is tied to Mohr–Coulomb slope-stability reference formulations, but the object studied here is the algorithmic endpoint surrogate defined in Section 5.5 and equations (22) and (32). The Sysala return-mapping, optimization, convex-optimization, and advanced continuation papers, together with the efficient Čermák–Sysala–Valdman MATLAB elastoplastic implementation of Čermák et al. [18], provide reference-model and method context for this problem class [2–5]; they are not used here as verified numeric paper baselines.

The endpoint-surrogate comparison therefore uses two diagnostics. The endpoint-observable check tests a seven-step direct load branch, $\lambda_{\text{sr}} = 1.0, 1.1, \dots, 1.6$, against the reference model on the $P_2(L_1)$ reference case. The work relative difference is 3.877×10^{-5} , the displacement relative L^2 difference is 3.517×10^{-3} , and the deviatoric-strain relative L^2 difference is 8.720×10^{-3} . The fixed-load reference comparison asks a narrower fixed- λ_{sr} question against a matched reference model; it is not proof of path-consistent incremental-history equivalence. After aligning the reference-model boundary conditions with the active boundary condition set, the fixed-load grid is capped at the main-figure load level $\lambda_{\text{sr}} = 1.55$. Here $\lambda_{\text{max}}^{\text{succ}}$ denotes the largest value in the tested grid $\{1.0, 1.2, 1.4, 1.5, 1.55\}$ that satisfies the relative-correction stopping criterion. The comparison matches $\lambda_{\text{max}}^{\text{succ}}$ on that grid and satisfies the predefined criteria for the $u_{\max}(\lambda_{\text{sr}})$ curve and endpoint displacement: the $u_{\max}(\lambda_{\text{sr}})$ relative L^2 difference is 6.396×10^{-6} , and the endpoint displacement relative L^2 difference is 8.299×10^{-4} . The endpoint deviatoric-strain relative L^2 difference (4.307×10^{-3}) and boundary-profile relative L^2 difference (5.229×10^{-4}) are reported as diagnostics rather than acceptance criteria. This comparison uses a matched mesh, material table, Davis-B strength reduction, load schedule,

boundary conditions, and active free degrees of freedom, but it remains a fixed-load endpoint check rather than a path-consistent plastic-history reference.



(a) Endpoint-observable comparison with the reference model.



(b) Fixed- λ_{sr} reference diagnostic.

Figure 20: Plasticity3D endpoint-surrogate comparison. The upper panel compares the endpoint observables after the seven-step direct branch ending at $\lambda_{\text{sr}} = 1.6$; the lower panel shows the fixed- λ_{sr} reference diagnostic after matching the boundary conditions and active free degrees of freedom.

Check	Comparison	Relative difference	Status
endpoint observable	work	3.88×10^{-5}	comparison
endpoint observable	displacement relative L^2	3.52×10^{-3}	comparison
endpoint observable	deviatoric-strain relative L^2	8.72×10^{-3}	comparison
fixed-load comparison	$\lambda_{\text{max}}^{\text{succ}}$	0	satisfied
fixed-load comparison	$u_{\text{max}}(\lambda_{\text{sr}})$ relative L^2	6.4×10^{-6}	satisfied
fixed-load comparison	endpoint displacement relative L^2	8.3×10^{-4}	satisfied
fixed-load comparison	endpoint deviatoric-strain relative L^2	4.31×10^{-3}	diagnostic
fixed-load comparison	boundary profile relative L^2	5.23×10^{-4}	diagnostic
fixed-load comparison	fixed-load criteria, 3/3 satisfied	–	satisfied

Table 14: Plasticity3D validation summary for the direct-branch endpoint-observable comparison ending at $\lambda_{\text{sr}} = 1.6$ and the fixed- λ_{sr} reference diagnostic.

The resulting interpretation is limited but useful. Plasticity3D is reported here through direct-branch endpoint-observable differences, while the fixed-load reference comparison supports matched endpoint observables for that surrogate under matched boundary conditions and active free degrees of freedom. Neither check is a path-consistent incremental-plasticity history validation.

7 Performance and Solver Behavior

This section turns from problem definition to computational behavior. The benchmark families in Section 5 established the governing equations, discrete energies, and representative states, and Section 6 separated the validation and external-comparison evidence. Here we focus on parallel timings, scaling trends, solver behavior, and the comparative cost of alternative derivative or discretization choices. Earlier timing evidence is retained when it clarifies scalability, but it is treated as performance evidence rather than validation evidence. The status and timing terms used below follow the protocol in Table 2.

Evidence block	Role	Solver policy	Stop or fixed work	Reported time
Scalar benchmarks	formulation, derivative-route, and globalization evidence	Newton with Armijo or the stated trust-region variant; CG or GMRES with Hypre in PETSc rows	configured nonlinear tolerances; capped rows are diagnostic	wall or solve/elapsed time in globalization rows
Hyperelasticity	finite-strain mechanics, JAX-FEM comparison, and scaling	trust-region or line-search Newton; GAMG or same-mesh PMG profiles as stated	per-load-step stationarity; fixed-step rows are solver diagnostics	wall time for benchmark rows; solver total for first-step scaling
Plasticity2D	endpoint evidence plus larger fixed-work diagnostics	Armijo continuation with same-mesh PMG smoother policy	$P_4(L_5)$ is the endpoint solve; $P_4(L_6)$ – $P_4(L_7)$ are fixed-iteration diagnostics	wall time for the stated endpoint or fixed-work run
Plasticity3D validation	endpoint-surrogate agreement, not path-history validation	fixed-load reference-model comparisons under matched boundary conditions	validation criteria use fixed-load observables; strain/profile rows are diagnostic	timing is secondary to agreement metrics
Plasticity3D performance	second-order routes, globalization, and parallel scaling	Armijo, residual-bisection, or trust-region Newton with FGMRES and same-mesh PMG	relative correction, gradient target, cap, or fixed work as stated	wall, solve, or solver-total time according to the table header
Topology	distributed design-mechanics timing and rank consistency	adaptive reduced-objective continuation with PETSc mechanics and GAMG	design/state-change stall criterion or fixed schedule	end-to-end wall time

Table 15: Numerical evidence blocks and reporting protocol used in this section. Exact tolerances and caps are stated with the corresponding result when they affect the interpretation.

7.1 Globalization Method Comparison

Table 16 isolates the nonlinear globalization choice on four benchmark configurations using the same element assembly path and fixed preconditioner family within each problem. The scalar cases use level-10 meshes obtained by one additional uniform refinement of the level-9 inputs. The p -Laplace globalization row is a separate unit-load ($f = -1$) stress test on the same L-shaped hierarchy, so its energies are not compared with the $f = -10$ benchmark and scaling rows in Sections 5.1 and 7.3. For p -Laplace, the comparison uses 32 MPI ranks with a 30 s wall-time cap. For Ginzburg–Landau, the reported comparison uses 16 MPI ranks with a 120 s wall-time cap. For hyperelasticity, the comparison uses the first eight load steps of the level-4 rank-local assembly path on 32 MPI ranks with a 270 s cap. The Plasticity3D comparison uses the $P_2(L_1)$ endpoint surrogate at $\lambda_{sr} = 1.55$ on 32 MPI ranks with a 180 s cap. Timed-out configurations are retained in the comparison; the table uses solver-reported solve time for completed solves and elapsed wall time for timed-out cases. In the line-search configurations the trust-region machinery is disabled; in particular, the hyperelasticity line-search baseline uses GMRES rather than the STCG trust-region subproblem solver. For the Plasticity3D configuration, all methods are judged against the stricter $\|\nabla\mathcal{F}\|_2 < 10^{-4}$ endpoint criterion.

The clean p -Laplace case is the expected easy case: Newton with Armijo line search is sufficient and gives the lowest solve time. Ginzburg–Landau separates the policies more sharply: the same line-search-only policy times out under the fixed cap, while both trust-region variants converge to the same reported energy. For the hyperelasticity excerpt all three policies complete, but Steihaug trust-region globalization without the line-search subproblem safeguard has the largest trust-region rejection count and a larger solve time than the hybrid policy, whereas line-search-only Newton is slowest. The hybrid policy keeps the trust-region safeguard while reducing both nonlinear iterations and solver time relative to line-search-only Newton. The Plasticity3D endpoint gives the opposite conclusion from the easy scalar case: at the stricter $\|\nabla\mathcal{F}\|_2 < 10^{-4}$ criterion, line-search-only Newton reaches the iteration cap, whereas both trust-region policies converge. The capped row matches the displayed energy but remains above the gradient target, as reported by the gradient-to-target column in Table 16. The hybrid variant is the fastest of the two trust-region configurations and also uses fewer Krylov iterations. Together, these comparisons support a problem-dependent policy choice rather than a universal preference for any single globalization strategy.

We additionally evaluate the Ginzburg–Landau level-10 case under an eight-rank budget. Table 17 reports this scalar-nonconvex comparison separately from the multi-family table above. Its purpose is narrower: it records that the trust-region variants complete with matching energy and Newton–Krylov work, while the line-search-only configuration still times out under the fixed wall-time cap. For that timed-out configuration we report elapsed wall time and the cap, whereas completed configurations report solver time.

7.2 Derivative-Route Comparison

Table 18 isolates the cost of the second-order route on three representative matched cases. The p -Laplace cases reuse the selected level-9, 32-rank timing data. The hyperelasticity cases run only the first level-4 load step and intentionally

Outcome and terminal value							
Benchmark	Method	Ranks	Outcome	Steps	Solve/elapsed [s]	Grad/target	Energy
p -Laplace L_{10}	Newton + LS	32	completed	1/1	6.95	–	–0.2517224496
p -Laplace L_{10}	Steihaug TR	32	completed	1/1	11.7	–	–0.2517227312
p -Laplace L_{10}	Hybrid TR+LS	32	completed	1/1	12	–	–0.2517227312
Ginzburg–Landau L_{10}	Newton + LS	16	timeout	0/1	121	–	–
Ginzburg–Landau L_{10}	Steihaug TR	16	completed	1/1	14.3	–	0.3456246783
Ginzburg–Landau L_{10}	Hybrid TR+LS	16	completed	1/1	14.3	–	0.3456246783
Hyperelasticity L_4	Newton + LS	32	completed	8/8	194	–	9.733145
Hyperelasticity L_4	Steihaug TR	32	completed	8/8	178	–	9.73101
Hyperelasticity L_4	Hybrid TR+LS	32	completed	8/8	149	–	9.730978
Plasticity3D $P_2(L_1)$	Newton + LS	32	iteration cap	0/1	70.2	1.18	–3,012,813
Plasticity3D $P_2(L_1)$	Steihaug TR	32	completed	1/1	20.8	0.00967	–3,012,813
Plasticity3D $P_2(L_1)$	Hybrid TR+LS	32	completed	1/1	16.4	0.579	–3,012,813
Nonlinear and Krylov work							
Benchmark	Method	Ranks	Newton	Krylov	LS evals	TR rejects	
p -Laplace L_{10}	Newton + LS	32	15	32	15	0	
p -Laplace L_{10}	Steihaug TR	32	29	32	0	0	
p -Laplace L_{10}	Hybrid TR+LS	32	29	32	29	0	
Ginzburg–Landau L_{10}	Newton + LS	16	0	0	0	0	
Ginzburg–Landau L_{10}	Steihaug TR	16	19	25	0	0	
Ginzburg–Landau L_{10}	Hybrid TR+LS	16	19	25	19	0	
Hyperelasticity L_4	Newton + LS	32	420	6984	2758	0	
Hyperelasticity L_4	Steihaug TR	32	314	7245	0	95	
Hyperelasticity L_4	Hybrid TR+LS	32	295	6849	499	3	
Plasticity3D $P_2(L_1)$	Newton + LS	32	80	858	224	0	
Plasticity3D $P_2(L_1)$	Steihaug TR	32	13	96	0	1	
Plasticity3D $P_2(L_1)$	Hybrid TR+LS	32	11	73	12	0	

Table 16: Fixed-case comparison of nonlinear globalization policies. LS denotes Armijo line search and TR denotes Steihaug trust region. The gradient-to-target column is reported for the Plasticity3D endpoint gate.

Method	Outcome	Newton	Krylov	LS evals	TR rejects	Solve/elapsed [s]	Cap [s]	Energy
Newton + LS	timeout	–	–	–	–	301	300	–
Steihaug TR	completed	19	27	0	0	28.5	–	0.3456246783
Hybrid TR+LS	completed	19	27	19	0	28.5	–	0.3456246783

Table 17: Ginzburg–Landau level-10 fixed-budget globalization comparison on eight MPI ranks. The configurations use the same element assembly and Hypre preconditioner; only the nonlinear globalization policy changes. The timed-out configuration reports elapsed wall time and the configured wall-time cap.

use the replicated element path for both routes, because the rank-local hyperelasticity configuration does not use colored SFD. This makes the case a derivative-route comparison rather than a scaling claim. The Plasticity3D cases use the same $P_2(L_1)$, $\lambda_{\text{sr}} = 1.55$, MUMPS-backed PMG setting as the globalization comparison above. Because the available timing records use solver timers, solver totals, and wall times in different cases, the timing-scope column states which quantity is used for each configuration.

Across all three benchmark cases, colored SFD reaches the same reported energy and identical nonlinear/Krylov iteration counts as element AD. The scalar and hyperelasticity cases still show the expected wall-time premium for colored SFD. On the Plasticity3D endpoint, constitutive AD is fastest, while element AD and colored SFD have matching iteration counts and comparable solve times. The colored-SFD Plasticity3D configuration records 240–420 colors across the 32 ranks; its comparable solve time is therefore explained by batched rank-local HVPs having nearly the same measured Hessian-callback cost as element AD on this $P_2(L_1)$ case, not by an under-colored finite-difference reconstruction.

Table 19 adds the degree sensitivity behind the Plasticity3D derivative-route claim. This single-linearization comparison uses one Newton linearization on bottom-clamped heterogeneous meshes at $\lambda_{\text{sr}} = 1.55$, 32 MPI ranks, and the degree-appropriate PMG hierarchy. This table should be read as derivative cost evidence, not as a convergence-quality replacement for the full Plasticity3D endpoint studies. The same-DOF comparison is the most informative

Outcome and timing

Benchmark	Route	Outcome	Time [s]	Timing scope	Energy
p -Laplace L_9	Element AD	completed	1.14	solver total	−7.960003
p -Laplace L_9	Colored SFD	completed	1.45	solver total	−7.960003
Hyperelasticity L_4 step 1	Element AD	completed	11.7	solve	0.1519923275
Hyperelasticity L_4 step 1	Colored SFD	completed	34.6	solve	0.1519923275
Plasticity3D $P_2(L_1)$					
$\lambda_{sr} = 1.55$	Element AD	completed	16.3	solve	−3,012,813
Plasticity3D $P_2(L_1)$					
$\lambda_{sr} = 1.55$	Colored SFD	completed	16.4	solve	−3,012,813
Plasticity3D $P_2(L_1)$					
$\lambda_{sr} = 1.55$	Constitutive AD	completed	4.81	solve	−3,012,813

Work and Hessian construction

Benchmark	Route	Ranks	Newton	Krylov	Hessian [s]	SFD colors
p -Laplace L_9	Element AD	32	6	11	–	–
p -Laplace L_9	Colored SFD	32	6	11	–	–
Hyperelasticity L_4 step 1	Element AD	32	22	492	–	–
Hyperelasticity L_4 step 1	Colored SFD	32	22	492	–	–
Plasticity3D $P_2(L_1)$						
$\lambda_{sr} = 1.55$	Element AD	32	11	73	6.57	–
Plasticity3D $P_2(L_1)$						
$\lambda_{sr} = 1.55$	Colored SFD	32	11	73	6.59	240–420
Plasticity3D $P_2(L_1)$						
$\lambda_{sr} = 1.55$	Constitutive AD	32	11	73	1.6	–

Table 18: Matched-case derivative-route comparison. Colored SFD denotes rank-local colored sparse recovery; the reported configurations use AD-HVP probes, while finite-difference gradient probes are the classical recovery variant. The SFD-color range is shown where per-rank color counts are available.

same-DOF contrast: $P_1(L_2)$ and $P_2(L_1)$ both have 79,024 free unknowns, but the P_2 element stencil is denser. Its colored-SFD route therefore needs 240–420 colors instead of 66–90, raises the Hessian callback from 0.332 to 1.412 s, and raises the rank-maximum RSS from about 0.65 to 1.10 GiB. The P_4 colored-SFD case is outside this single-linearization table; the P_2 same-DOF comparison already isolates the stencil-density effect.

Discretization size

Element	Route	Work	Free DOFs	Elem.	Overlap	RSS [GiB]
$P_1(L_1)$	Element AD	one Newton step	10,526	12	675	0.5
$P_1(L_1)$	Colored SFD	one Newton step	10,526	12	675	0.5
$P_1(L_1)$	Constitutive AD	one Newton step	10,526	12	675	0.5
$P_1(L_2)$	Element AD	one Newton step	79,024	12	3624	0.6
$P_1(L_2)$	Colored SFD	one Newton step	79,024	12	3624	0.6
$P_1(L_2)$	Constitutive AD	one Newton step	79,024	12	3624	0.5
$P_2(L_1)$	Element AD	one Newton step	79,024	30	3933	0.9
$P_2(L_1)$	Colored SFD	one Newton step	79,024	30	3933	1.1
$P_2(L_1)$	Constitutive AD	one Newton step	79,024	30	3933	0.6
$P_4(L_1)$	Element AD	one Newton step	610,964	105	27177	1.2
$P_4(L_1)$	Constitutive AD	one Newton step	610,964	105	27177	1.1

Linearization cost

Element	Route	Free DOFs	Krylov	Solve [s]	Hess. [s]	Colors
$P_1(L_1)$	Element AD	10,526	11	0.0806	0.257	–
$P_1(L_1)$	Colored SFD	10,526	11	0.0704	0.256	57–90
$P_1(L_1)$	Constitutive AD	10,526	11	0.0706	0.217	–
$P_1(L_2)$	Element AD	79,024	3	0.232	0.322	–
$P_1(L_2)$	Colored SFD	79,024	3	0.273	0.332	66–90
$P_1(L_2)$	Constitutive AD	79,024	3	0.164	0.257	–
$P_2(L_1)$	Element AD	79,024	5	0.83	1.31	–
$P_2(L_1)$	Colored SFD	79,024	5	0.827	1.41	240–420
$P_2(L_1)$	Constitutive AD	79,024	5	0.304	0.535	–
$P_4(L_1)$	Element AD	610,964	13	8.48	7.04	–
$P_4(L_1)$	Constitutive AD	610,964	13	3.97	3.62	–

Table 19: Plasticity3D one-Newton derivative-route cost versus discretization. Elem DOFs is the local vector element size used by the Hessian callback; overlap DOFs is the maximum local overlap vector size recorded during assembly.

7.3 p -Laplace

The p -Laplace problem definition and discrete energy are given in Section 5.1 and equation (7). In this clean scalar setting, the numerical story is easiest to interpret: exact element Hessians remain effective in the reported distributed runs, while colored SFD reaches nearly identical discrete energies at higher wall time.




Figure 21: p -Laplace strong scaling on the finest tested grid. All curves correspond to the same discrete problem from equation (7).

In the finest plotted case, the JAX+PETSc element-AD path reports 1.14s on 32 ranks, while colored SFD reports 1.45s with the same Newton and Krylov iteration counts. The final energies agree to the displayed precision, so this block mainly measures the overhead of sparse recovery on a smooth scalar problem. Across the plotted ranks, element AD drops from 10.13s on one rank to 1.14s on 32 ranks; colored SFD follows the same trend with persistent recovery overhead.

7.4 Ginzburg–Landau

The Ginzburg–Landau family from Section 5.2 and equation (11) shows the same high-level derivative story under a more indefinite linearization. The important computational message is not a dramatic separation between exact and approximate second-order routes, but that the distributed solver remains convergent on this nonconvex scalar energy with the same sparse infrastructure.




Figure 22: Ginzburg–Landau strong scaling for the tested fine-grid case governed by equation (11).

At 32 ranks, the element-AD and colored-SFD JAX+PETSc runs report 2.10s and 2.72s, respectively, and both reach the same displayed energy with identical nonlinear and Krylov counts. This supports the same derivative-route

interpretation as in the p -Laplace case, but on a nonconvex scalar energy. Across the plotted ranks, element AD falls from 15.64s on one rank to 2.10s on 32 ranks, with a nonideal plateau around 4–8 ranks before the 32-rank point.

7.5 Hyperelasticity

The hyperelasticity problem definition, boundary data, and terminal energies are given in Section 5.3 and equation (15). Computationally, this is the benchmark in which globalization matters more than derivative-route choice: the distributed implementation completes the full load path and preserves the terminal energy trend from Section 5.3. The external JAX-FEM comparison is reported separately in Section 6.2; the figure below reports only distributed scaling for the load path in equation (15).




Figure 23: Hyperelasticity strong scaling for the 24-step load path defined by equation (15).

The 24-step hyperelasticity scaling curve reduces JAX+PETSc element-AD wall time from 1894s at 8 ranks to 529s at 32 ranks. The speedup is substantial but not ideal, which motivates the more targeted fixed-work distribution and PMG diagnostics below.

The larger multi-node experiment fixes the nonlinear work to the first level-5 load step after the rank-local hyperelasticity implementation. The CPU nodes used here have two AMD EPYC 7H12 64-core processors and 256 GB RAM [38]; the 128 cores are arranged as eight NUMA domains with 16 cores each. The runs use one MPI rank per core with contiguous CPU binding and NUMA-local memory binding. The mesh is generated procedurally from the uniform beam hierarchy, and the element path assembles from rank-local elements with rank-local PETSc COO preallocation and point-to-point overlap exchange.

The nonlinear solve is the same trust-region Newton path used for the benchmark in section 5.3. The linear subproblems use STCG with a same-hierarchy PMG preconditioner, Chebyshev/Jacobi smoothing, and coarsest level L_2 . The coarse solve is PETSc redundant LU with MUMPS: the partial one-node configurations use one redundant group over the active ranks, while the multi-node configurations use one 128-rank redundant group per node. Figure 24 and table 22 report measured solver total time rather than scheduler elapsed time. All configurations reach the same nonlinear and linear iteration counts shown in Table 22: 21 Newton iterations and 43 Krylov iterations.

Table 20 reports the fixed-cost memory comparison that motivated the rank-local assembly path used in the multi-node runs. The same-work check compares the replicated and rank-local level-4 first-step solve at four MPI ranks; the memory-scaling cases use one Newton linearization on the level-5 mesh at 8, 16, and 32 ranks. The table reports energy, operating-system peak RSS, and the measured footprint of assembly arrays, making clear whether a configuration stores global mesh and adjacency data or only local elements plus overlap.

The level-4 same-work rows show matching Newton and Krylov counts and matching energy to the displayed precision for replicated and rank-local assembly. On the larger level-5 one-linearization memory rows, rank-local assembly reduces summed RSS from 215.6 GiB to 104.9 GiB at 32 ranks while leaving solve time at the same reported precision.

Table 21 isolates the first-step level-5 PMG choice on the same 32-rank level-5 problem with fixed nonlinear work. All configurations perform the same eight Newton linearizations, so the table compares the cost and terminal state reached under different preconditioner and coarse-solver policies. The energy column is the state after this fixed nonlinear work, so Table 21 is a solver-policy comparison rather than an equal-accuracy timing comparison or endpoint-accuracy ranking.

<i>Solve work</i>								
Purpose	Assembly layout	Outcome	Level	Ranks	Newton	Krylov	Solve [s]	Energy
same-work check	replicated	completed	L_4	4	26	571	51.2	0.151971
same-work check	rank-local	completed	L_4	4	26	571	50.8	0.151971
memory comparison	replicated	one linearization	L_5	8	1	1	10.7	53.95804
memory comparison	rank-local	one linearization	L_5	8	1	1	10.5	53.95926
memory comparison	replicated	one linearization	L_5	16	1	1	5.57	32.927343
memory comparison	rank-local	one linearization	L_5	16	1	1	5.5	32.927867
memory comparison	replicated	one linearization	L_5	32	1	1	3.37	23.942407
memory comparison	rank-local	one linearization	L_5	32	1	1	3.36	23.803027

<i>Memory and overlap</i>						
Purpose	Assembly layout	Level	Ranks	RSS max/sum [GiB]	Tracked sum [GiB]	Overlap/owned
same-work check	replicated	L_4	4	3.6/14.2	5.49	1.01
same-work check	rank-local	L_4	4	3.3/13.2	6.01	1.01
memory comparison	replicated	L_5	8	14.9/118.7	44.34	1.01
memory comparison	rank-local	L_5	8	12.3/98.3	47.87	1.01
memory comparison	replicated	L_5	16	9.4/150.9	45.45	1.03
memory comparison	rank-local	L_5	16	6.3/100.3	47.89	1.03
memory comparison	replicated	L_5	32	6.7/215.6	47.65	1.05
memory comparison	rank-local	L_5	32	3.3/104.9	47.92	1.05

Table 20: Rank-local hyperelasticity distribution and memory comparison. RSS max/sum gives the maximum and summed per-rank resident set size. Tracked sum is the reported assembly-array footprint, and overlap/owned is the summed rank-local overlap DOF count divided by global owned free DOFs.

<i>Outcome and solver work</i>						
Precond.	Nonlinear work		Newton	Krylov	Solver [s]	Energy
GAMG	8 Newton linearizations		8	42	22.3	1.608441
PMG L_2 + Hypre	8 Newton linearizations		8	92	23.3	0.5252349427
PMG L_2 + MUMPS	8 Newton linearizations		8	17	15.3	0.3848912708
PMG L_3 + MUMPS	8 Newton linearizations		8	17	16.3	0.3864166034

<i>Time breakdown and coarse solve</i>					
Precond.	Coarse	Assembly [s]	PC [s]	Linear [s]	
GAMG	–	9.11	6.41	3.91	
PMG L_2 + Hypre	Hypre	10.3	2.93	8.14	
PMG L_2 + MUMPS	MUMPS, one redundant group	9.13	2.55	1.91	
PMG L_3 + MUMPS	MUMPS, one redundant group	9.14	2.93	2.51	

Table 21: Rank-local hyperelasticity level-5 PMG and coarse-solver sensitivity at 32 MPI ranks with fixed nonlinear work. The PMG configurations use the same rank-local element path as the scaling runs.

The fixed-work table shows that the MUMPS-coarse PMG variants reduce total Krylov iterations from 92 with Hypre coarse solves to 17, and reduce solver time from 23.3s to about 16s. This is preconditioner-policy evidence for a fixed nonlinear workload, not a statement about final load-path accuracy.




Figure 24: Multi-node CPU level-5 hyperelasticity first-step scaling at fixed nonlinear work with the PMG coarse level fixed at L_2 . Multi-node annotations show node counts.

Ranks	Nodes	Coarse groups	Ranks/group	Solver total [s]	First step [s]	Newton iters	Krylov iters
16	1	1	16	419	367	21	43
32	1	1	32	317	280	21	43
64	1	1	64	148	132	21	43
128	1	1	128	67.4	58.4	21	43
256	2	2	128	36.4	30.7	21	43
512	4	4	128	21.7	17.7	21	43
1024	8	8	128	15.7	12.3	21	43

Table 22: Multi-node CPU level-5 hyperelasticity first-step timings used in figure 24.

Figure 24 and table 22 show a fixed-work timing drop from 419s on 16 ranks to 15.7s on 1024 ranks, with the same 21 Newton iterations and 43 Krylov iterations in all listed configurations. This evidence isolates the multi-node sparse-solver and PMG-coarse-level behavior for one load step, rather than the accuracy of the full load path.

7.6 Plasticity2D

The Plasticity2D model and Davis-B reduction are defined in Section 5.4 and equations (18) and (19). For this family, the most informative computational evidence in the present paper is not a single strong-scaling curve. The benchmark section separates the completed endpoint case from the larger fixed-cost studies in Figure 13 and table 10, and the reference-continuation comparison is collected later in Section A. Thus the $P_4(L_5)$ case is the endpoint evidence for this family, while the $P_4(L_6)$ and $P_4(L_7)$ cases are fixed-work diagnostics for nonlinear-tail and PMG behavior. The 8-rank and 32-rank reference-continuation rows in Table 31 support solver-policy timing under a common continuation logic; they are not a separate validation result or a strong-scaling claim.

7.7 Plasticity3D

The governing problem, Davis-B reduction, scalar constitutive potential, and deviatoric-strain norm are defined in Section 5.5 and equations (22), (24), (32) and (33). The validation limits for this endpoint surrogate are reported in Section 6.3. With that interpretation fixed, this section asks two performance questions: which second-order route is fastest on the fixed high-order derivative-route case under matched terminal observables, and how the endpoint surrogate behaves across degree, resolution, and rank count.

Derivative-route comparison. The first question is which second-order route gives the lowest measured cost without changing the reported terminal observables. On the fixed $P_4(L_1)$ case at $\lambda_{sr} = 1.5$ and 8 MPI ranks, all three routes converge to the same reported terminal observables, with identical Newton and total Krylov counts. This residual-bisection Newton comparison uses the HyPre-backed rank-local linear-solver profile, so it should be read as derivative-route timing evidence rather than as the MUMPS-backed PMG profile used by the later $\lambda_{sr} = 1.55$ convergence studies. The differences are cost differences, not state-quality differences: constitutive AD has the lowest median wall time on this fixed case, with median wall time 222s and median solve time 191s, compared with 288s/255s for element AD and 372s/253s for colored SFD. In Table 23, ω denotes the external-work observable and u_{max} denotes the maximum displacement magnitude; these columns check that the derivative route does not change the reported terminal observables.





Figure 25: Fixed Plasticity3D derivative-route comparison on $P_4(L_1)$, $\lambda_{sr} = 1.5$, and 8 MPI ranks.

Route	Wall time [s]	Solve time [s]	Newton iters	Krylov iters
Element AD	288	255	8	376
Constitutive AD	222	191	8	376
Colored SFD	372	253	8	376


Terminal observables				
Route	Energy	ω	$u_{\max}$	
Element AD	-3,010,861	6,027,722	0.731848	
Constitutive AD	-3,010,861	6,027,722	0.731848	
Colored SFD	-3,010,861	6,027,722	0.731848	

Table 23: Plasticity3D derivative-route comparison on the fixed $P_4(L_1)$, $\lambda_{\text{sr}} = 1.5$, 8-rank case. These configurations use the HyPre-backed rank-local profile and are separate from the MUMPS-backed PMG configurations in the $\lambda_{\text{sr}} = 1.55$ studies.


Discretization and scaling. The remaining figures in this subsection answer a different question: how the endpoint surrogate behaves across degree/resolution choices, and how the selected solver configurations scale on the separate $P_4(L_2)$ timing studies. The zoomed panels in Figure 26 are included because the full figure is dominated by coarse P_1 variation. Under the bottom-clamped constraint, the matched set at 610,964 free DOFs is $P_1(L_3)$, $P_2(L_2)$, and $P_4(L_1)$. The matched set at 4,801,816 free DOFs is $P_1(L_4)$, $P_2(L_3)$, and $P_4(L_2)$. The scientific message is therefore that the solver reaches similar terminal energies along both p - and h -refined paths, while higher order on coarser meshes buys lower observed values of equation (22) at matched DOFs only by paying a clear wall-time premium in this nine-case endpoint study.




(a) All degrees: final energy versus free DOFs.



(b) All degrees: final energy versus total wall time.



(c) P_2/P_4 zoom: final energy versus free DOFs.



(d) P_2/P_4 zoom: final energy versus total wall time.

Figure 26: Bottom-clamped Plasticity3D degree-versus-resolution study at $\lambda_{\text{sr}} = 1.55$. The top panels show the full $P_1/P_2/P_4$ comparison, while the bottom panels zoom to P_2/P_4 so the near-overlap hidden by the coarse P_1 points becomes visible. The left column plots final values of equation (22) against free DOFs; the right column plots the same converged energies against end-to-end wall time. All timings use the same 32-rank configuration.




Figure 27: Auxiliary strong scaling for the selected Plasticity3D solver configuration: constitutive-AD rank-local assembly with same-mesh PMG on $P_4(L_2)$ at $\lambda_{sr} = 1.0$. This is a separate timing study from the bottom-clamped $\lambda_{sr} = 1.55$ results shown in figures 14, 15 and 26.

Ranks	Wall time [s]	Solve time [s]	Speedup	Efficiency	Newton iters	Krylov iters
1	5424	5176	1	1	17	114
2	2914	2777	1.861	0.931	17	114
4	1717	1634	3.16	0.79	17	114
8	1164	1107	4.658	0.582	17	114
16	608	575	8.916	0.557	17	114
32	300	284	18.053	0.564	17	114

Table 24: Auxiliary Plasticity3D strong scaling for constitutive-AD rank-local assembly with same-mesh PMG on $P_4(L_2)$ at $\lambda_{sr} = 1.0$. These timings provide solver evidence for a separate load factor from the bottom-clamped $\lambda_{sr} = 1.55$ benchmark figures.

Figure 27 and table 24 give a high-DOF Plasticity3D distributed timing result for the Hypre-coarse auxiliary PMG profile. This scaling evidence uses the selected $P_4(L_2)$ mesh family but a different load factor, $\lambda_{sr} = 1.0$, so it complements rather than extends the bottom-clamped $\lambda_{sr} = 1.55$ benchmark figures above. Within that timing study, the reported nonlinear and linear iteration counts remain matched on 1/2/4/8/16/32 ranks, while the end-to-end time drops from 5424s on one rank to 300s on 32 ranks. Additional reference-formula assembly and continuation evidence is reported later in Section A.

The converged bottom-clamped $P_4(L_2)$, $\lambda_{sr} = 1.55$ scaling run uses the same case and load level as the benchmark table, with the constitutive-AD rank-local assembly and MUMPS-backed PMG profile. The PMG hierarchy is the same-mesh $P_4 \rightarrow P_2 \rightarrow P_1$ hierarchy, the nonlinear method is Armijo Newton, and the PETSc relative KSP residual tolerance is 1×10^{-2} . The nonlinear stop requires both relative correction below 0.002 and a stopping-gradient norm below one percent of the first Newton gradient. Figure 28 and table 25 compare the single-node CPU sweep with the multi-node CPU sweep at 16 MPI ranks per node. All plotted configurations converge in 29 Newton iterations and 1240 total Krylov iterations to the same energy and stopping-gradient norm to the precision shown in the table. The apparently large absolute value is below the run-specific relative gradient target; the table therefore reports both $\|g\|_{\text{stop}}$ and the ratio $\|g\|_{\text{stop}}/\|g\|_{\text{target}}$. The timing comparison is not confounded by different nonlinear work.




Figure 28: Single-node and multi-node CPU converged scaling for the bottom-clamped Plasticity3D $P_4(L_2)$ case at $\lambda_{sr} = 1.55$. Dashed references are ideal $1/r$ lines anchored separately to the first point of each CPU setting.

CPU setting	Ranks	Nodes	Solver total [s]	Solve [s]	Newton iters	Krylov iters	Energy	$\ g\ _{stop}$	$\ g\ _{stop}/target$
single-node CPU	1	—	17,890	17,615	29	1240	−3,013,348	10.594	0.525
single-node CPU	2	—	9574	9412	29	1240	−3,013,348	10.594	0.525
single-node CPU	4	—	5667	5564	29	1240	−3,013,348	10.594	0.525
single-node CPU	8	—	4003	3947	29	1240	−3,013,348	10.594	0.525
single-node CPU	16	—	2022	1993	29	1240	−3,013,348	10.594	0.525
single-node CPU	32	—	1110	1093	29	1240	−3,013,348	10.594	0.525
multi-node CPU	16	1	2246	2204	29	1240	−3,013,348	10.594	0.525
multi-node CPU	32	2	1251	1227	29	1240	−3,013,348	10.594	0.525
multi-node CPU	64	4	769	754	29	1240	−3,013,348	10.594	0.525
multi-node CPU	128	8	498	488	29	1240	−3,013,348	10.594	0.525

Table 25: Converged timings for the Plasticity3D $P_4(L_2)$, $\lambda_{sr} = 1.55$ run in figure 28. Single-node configurations use one CPU node; multi-node configurations use 16 ranks per CPU node. The nonlinear stop requires relative correction below 0.002 and final gradient below one percent of the first Newton gradient. The stopping-gradient norm in this table is the relative-gradient stopping metric for this scaling run, not the absolute final-gradient column in the degree benchmark table.

The distributed layout measurements in Table 26 identify one observed limit on ideal scaling. The maximum local overlap vector shrinks with rank count, but the sum of rank-local overlap DOFs increases from about the owned problem size on one rank to more than four times the owned size on 128 ranks. This overlap replication comes from subdomain boundaries and high-order element support, so it is a decomposition cost rather than duplicated global mesh storage.

CPU setting	Ranks	Nodes	Max local DOFs	Overlap / owned DOFs
single-node CPU	1	—	4,859,595	1.01
single-node CPU	2	—	2,495,265	1.05
single-node CPU	4	—	1,284,351	1.11
single-node CPU	8	—	669,783	1.21
single-node CPU	16	—	351,219	1.41
single-node CPU	32	—	182,613	1.83
multi-node CPU	16	1	351,219	1.41
multi-node CPU	32	2	182,613	1.83
multi-node CPU	64	4	91,821	2.66
multi-node CPU	128	8	50,637	4.31

Table 26: Distributed layout measurements for the converged Plasticity3D scaling run. “Max local DOFs” is the largest local-overlap vector size on any rank; “Overlap / owned DOFs” is the summed rank-local overlap DOF count divided by the global owned free-DOF count.

7.8 Topology Optimization

The reduced design problem, mechanics subproblem, and objective history are defined in Section 5.6 and equations (35) and (36). The numerical question here is distributed outer-loop stability: the mechanics solve and the design update must remain synchronized under continuation in p and move-regularized updates. Because the outer-loop work can vary with the continuation schedule, the timing evidence here should be read as end-to-end timing under adaptive stopping rather than strong scaling at a fixed computational workload.




Figure 29: Topology end-to-end parallel timing trend for the fine-grid run. The plotted wall times correspond to the coupled reduced-design problem defined by equation (36), with per-rank stopping behavior reported in table 27.

Ranks	Wall time [s]	Solve time [s]	Outer iters	p	Compliance	Volume fraction
1	208	200	65	4.2	9.1559	0.3884
2	138	134	72	5.6	8.9473	0.3932
4	95.7	92.8	69	4.6	9.1684	0.385
8	82.1	80.1	67	7	9.2178	0.3932
16	43.3	41.9	66	5.8	9.6859	0.3798
32	25.1	23.7	66	6.6	9.9074	0.3749

Table 27: Parallel topology summary for the fine-grid study, reported for the reduced objective in equation (36).

Figure 29 and table 27 show that the fine-grid end-to-end wall time falls from 208 s on one rank to 25.1 s on 32 ranks. The outer iteration counts and final continuation exponents vary across ranks, so this is adaptive-stopping timing evidence for the coupled design-mechanics computation rather than a fixed-work strong-scaling claim.

Table 28 complements the end-to-end timing table with a controlled 40-iteration rank-consistency test on a smaller 384×192 grid. Early convergence and stall stopping are disabled, and all rank counts use the same initial design and smoother continuation schedule ($p \leftarrow p + 0.1$). The final density field is compared with the one-rank reference after 40 outer iterations. All rank counts complete the same 40 outer iterations. The compliance difference stays below 0.2%, and the density relative L^2 difference remains below 0.04, which checks that the distributed design-gradient exchange and mechanics coupling remain bounded under rank variation at this resolution.

Ranks	Schedule	Outer	Solve [s]	Compliance	Volume	p	$\Delta C/C_1$	Density rel. L^2
1	fixed schedule	40	48.3	15.56	0.3085	5	0	0
2	fixed schedule	40	28.5	15.56	0.3083	5	1.29×10^{-6}	1.37×10^{-2}
4	fixed schedule	40	19.7	15.5308	0.3086	5	1.87×10^{-3}	1.91×10^{-2}
8	fixed schedule	40	16	15.5665	0.3081	5	4.16×10^{-4}	3.68×10^{-2}
16	fixed schedule	40	8.23	15.5606	0.3084	5	4.15×10^{-5}	1.47×10^{-2}
32	fixed schedule	40	6.43	15.5637	0.3083	5	2.41×10^{-4}	9.71×10^{-3}

Table 28: Topology controlled 40-iteration rank-consistency check. Each rank count completes the same 40 outer iterations; compliance and density differences are relative to the one-rank case.

The topology benchmark completes the paper’s computational picture. Here the primary scalable JAX+PETSc realization is less about exact Hessians than about stable distributed design-mechanics coupling, while pure JAX remains the compact serial reference path discussed earlier in Table 5.

7.9 Synthesis of the Computational Evidence

The results above should be read as a coupled solver study rather than as a collection of isolated benchmark timings. First, the globalization tables show that the nonlinear policy is problem-dependent: line-search Newton is adequate on the clean p -Laplace case, while trust-region safeguards are beneficial on the nonconvex Ginzburg–Landau tests and needed for convergence in the stricter $P_2(L_1)$ Plasticity3D globalization test reported here. Second, the derivative-route comparisons show that access to several second-order routes is useful in practice. Colored SFD is a comparison and fallback route on the tested fixed-size problems, while constitutive AD gives the lowest measured wall time on the fixed high-order Plasticity3D derivative-route case. Third, the hyperelasticity distribution and PMG comparisons, together with the Plasticity3D partitioning measurements, show that sparse ownership, overlap size, coarse solves, and preconditioner policy are part of the numerical method, not merely software detail. Finally, the hyperelasticity, Plasticity3D, and topology scaling runs show that the primary JAX+PETSc realization solves the largest tested cases in this study, while serial reference formulations mainly serve validation or scoped comparison roles.

This synthesis is also the boundary of the paper’s claim. The experiments support a nonlinear FEM toolset in which local differentiable models, distributed sparse assembly, nonlinear globalization, Krylov linear algebra, and problem-specific preconditioners are developed and assessed together. They do not imply that one backend or one globalization policy dominates all finite-element computations, nor do they turn the Plasticity3D endpoint surrogate into a path-consistent incremental-history validation.

8 Discussion

The numerical evidence supports a single methodological lesson for the tested benchmark classes: local differentiation and sparse solver policy are coupled numerical choices. In several fixed-case comparisons, derivative routes and ownership layouts preserve displayed terminal observables while changing cost; globalization and preconditioner policies also change robustness, work, and timing. The following paragraphs delimit that lesson by derivative route, validation scope, and external comparison evidence.

Derivative routes interact with the sparse solver. Figure 25 and table 23 show a consistent cost comparison on the fixed $P_4(L_1)$, $\lambda_{\text{sr}} = 1.5$, 8-rank case. The three routes report terminal observables that agree to the displayed precision, with 8 Newton iterations and 376 Krylov iterations, so the comparison is about derivative cost rather than about different converged solutions. In this fixed test, constitutive AD preserves the same scalar-energy definition and has the lowest measured wall time: 222s, compared with 288s for element AD and 372s for colored SFD.

Three-dimensional plasticity validation is endpoint-scoped. The three-dimensional plasticity endpoint-surrogate comparison in Figure 20 and table 14 supports a scoped mechanics interpretation. The reported endpoint observables satisfy the stated fixed-load tolerances under the matched boundary-condition setup. The endpoint displacement is part of that fixed-load criterion, whereas the deviatoric-strain and boundary-profile comparisons are diagnostic. The evidence supports final-load scalar and field observables for the tested endpoint surrogate, but it does not validate incremental plastic-history paths or establish a numeric reproduction of the cited incremental-history studies.

External comparisons are useful when definitions are matched. The hyperelastic JAX-FEM comparison in Figure 19 and table 13 is informative because the mesh, prescribed displacement schedule, and terminal-state energy evaluation are explicitly matched. The terminal states are post-compared with the energy in equation (12); the final energy relative difference is 8.244×10^{-6} , and the displacement and centerline relative differences are below the prescribed 5% threshold. Its role is solver-consistency evidence for this hyperelastic benchmark: under a matched mesh, displacement schedule, and terminal-energy post-comparison, the JAX+PETSc and JAX-FEM terminal states agree closely.

Scope of the evidence. The same evidence also defines the scope of the conclusions. The $\lambda_{\text{sr}} = 1.0$ three-dimensional plasticity strong-scaling study is timing-only evidence, while the bottom-clamped $\lambda_{\text{sr}} = 1.55$ scaling sweep is the converged comparison tied to the benchmark. The topology study reports end-to-end timing under adaptive stopping, with a separate controlled rank-consistency check. The fixed-reference PMG studies should be read as solver-sensitivity tests, not external baselines. Together, these distinctions state the contribution: a nonlinear FEM toolset that couples JAX-based local differentiation with PETSc sparse distributed solvers and reports matched comparison evidence, with explicit agreement limits where predefined. Automatic differentiation alone does not determine the behavior of large nonlinear FEM; nonlinear globalization, MPI assembly, coarse solvers, and preconditioning also enter the numerical method.

9 Conclusion

We developed a JAX+PETSc toolset for nonlinear finite-element energy minimization in which local differentiable energy models are coupled to PETSc sparse distributed nonlinear solvers. The experiments connect scalar energy definitions, second-order derivative construction, MPI assembly, and reference comparisons in one solver setting. Across the tested problems, the main lesson is that local AD choices and distributed nonlinear solver design must be considered together.

The numerical results support this pattern on a deliberately heterogeneous benchmark suite. Scalar problems check formulation and derivative agreement, hyperelasticity provides a narrow matched external comparison, three-dimensional plasticity tests nonsmooth constitutive structure and distributed sparse solves under an endpoint-surrogate interpretation, and topology optimization exercises a coupled design-and-mechanics problem. In the fixed high-order three-dimensional plasticity derivative-route comparison, constitutive AD is the lowest measured wall-time route among the tested alternatives. In the reported $P_4(L_2)$, $\lambda_{sr} = 1.55$ three-dimensional plasticity scaling sweep, all rank counts converge with the same nonlinear and Krylov work. The topology runs provide distributed timing evidence for the coupled design-and-mechanics loop, with a separate controlled rank-consistency check.

Natural extensions are path-dependent mechanics validation against incremental-history references, broader agreement-focused external comparisons under matched problem definitions, and reducing preconditioner-setup and multigrid costs in the largest sparse mechanics runs in this study. Together, these directions extend the validation discipline of the present study toward broader nonlinear multiphysics and design problems.

A Solver-Policy Diagnostics

This appendix collects supporting solver-policy diagnostics that are methodologically useful but narrower than the main benchmark narratives.

A.1 Plasticity3D Reference-Formula and Fixed-Reference PMG Comparisons

The primary Plasticity3D numerical story in Section 7.7 is the rank-local constitutive-AD path. The tables below add two solver-policy comparisons on the same $P_4(L_2)$ problem from equation (22) at $\lambda_{sr} = 1.0$: a matched reference-formula assembly comparator and a fixed-reference PMG variant. Both use Armijo Newton, a gradient-only stopping target $\|\nabla\mathcal{F}\| < 0.01$, at most 50 Newton iterations, and PETSc relative KSP tolerance 1×10^{-1} . Here reference-formula assembly means closed-form constitutive assembly used as a comparator for the JAX constitutive-AD kernel, while fixed-reference PMG means a preconditioning variant in which the multigrid operator is not rebuilt along the nonlinear path.

Ranks	AD branch wall [s]	Closed-form branch wall [s]	AD branch solve [s]	Closed-form branch solve [s]	Ratio
4	1717	951	1634	890	1.805
8	1164	697	1107	656	1.671
16	608	374	575	350	1.628
32	300	277	284	260	1.085

Table 29: Matched constitutive-assembly versus reference-formula assembly comparison for the same PMG solver configuration on $P_4(L_2)$ at $\lambda_{sr} = 1.0$. The table reports wall time, solve time, and their ratio for matched rank counts on equation (22).

These diagnostics support two narrower statements. First, reference-formula assembly has lower reported wall time than rank-local assembly in the listed matched-rank comparisons, although the gap narrows at higher rank counts. Second, the fixed-reference PMG profile is not the default used in the present study: it may appear favorable on completed low-rank configurations, but it reaches the maximum number of Newton iterations on more of the reported rank counts and therefore gives less consistent convergence evidence. The rank-local constitutive-AD path remains the primary Plasticity3D realization because it uses the same local differentiable formulation as the derivative-route and converged scaling studies; the reference-formula configurations are retained to expose solver-policy sensitivity on matched discrete problems.

Ranks	PMG operator and tangent	Status	Wall time [s]
4	frozen PMG operator, AD branch tangent	completed	2116
4	frozen PMG operator, closed-form branch tangent	completed	988
8	frozen PMG operator, AD branch tangent	completed	1874
8	frozen PMG operator, closed-form branch tangent	completed	898
16	frozen PMG operator, AD branch tangent	iteration cap	1304
16	frozen PMG operator, closed-form branch tangent	iteration cap	601
32	frozen PMG operator, AD branch tangent	completed	348
32	frozen PMG operator, closed-form branch tangent	iteration cap	398

Newton–Krylov work				
Ranks	PMG operator and tangent	Newton iters	Krylov iters	Final gradient norm
4	frozen PMG operator, AD branch tangent	24	204	6.53×10^{-3}
4	frozen PMG operator, closed-form branch tangent	21	210	2.49×10^{-3}
8	frozen PMG operator, AD branch tangent	33	238	2.53×10^{-3}
8	frozen PMG operator, closed-form branch tangent	27	246	9.53×10^{-3}
16	frozen PMG operator, AD branch tangent	50	196	2.21×10^2
16	frozen PMG operator, closed-form branch tangent	50	177	7.35×10^2
32	frozen PMG operator, AD branch tangent	29	250	2.24×10^{-3}
32	frozen PMG operator, closed-form branch tangent	50	193	8.25×10^1

Table 30: Fixed-reference PMG alternative on the same $P_4(L_2)$, $\lambda_{sr} = 1.0$ Plasticity3D problem. The status column distinguishes completed configurations from rank counts that reach the maximum number of Newton iterations before satisfying the target gradient tolerance.

A.2 Plasticity2D Reference Continuation Evidence

For Plasticity2D, the reference-continuation comparison is not the primary method of the paper, but it is still useful as fixed-policy rank and timing evidence for the PMG smoother policy under the same outer continuation logic.

Policy	Ranks	Runtime [s]	Init Krylov iters	Continuation Krylov iters	Final λ_{sr}
fixed PMG smoother policy	8	489	58	733	1.568403
fixed PMG smoother policy	32	198	60	829	1.568334

Table 31: Reference continuation with the fixed PMG smoother policy at 8 and 32 ranks. This is supporting solver evidence for a reference continuation setting, not a primary solver comparison.

The two rows show fixed-policy timing evidence rather than a validation baseline: increasing from 8 to 32 ranks reduces the reported wall time from 489s to 198s, while the final strength-reduction values remain close at 1.568403 and 1.568334.

Code and Data Availability

The source code, input data, processed benchmark summaries, and figure and table generation scripts supporting this study are available at https://github.com/Beremi/fenics_nonlinear_energies. No separate archival DOI is cited for this version.

References

- [1] Franz Tschuchnigg, Helmut Schweiger, and Scott W. Sloan. Slope stability analysis by means of finite element limit analysis and finite element strength reduction techniques. part i: Numerical studies considering non-associated plasticity. *Computers and Geotechnics*, 70:169–177, 2015. doi: 10.1016/j.compgeo.2015.06.018. URL <https://tugraz.elsevierpure.com/en/publications/slope-stability-analysis-by-means-of-finite-element-limit-analysis>.
- [2] Stanislav Sysala, Martin Čermák, and Tomáš Ligurský. Subdifferential-based implicit return-mapping operators in Mohr-Coulomb plasticity. *ZAMM – Journal of Applied Mathematics and Mechanics / Zeitschrift für Angewandte Mathematik und Mechanik*, 97(12):1502–1523, 2017. doi: 10.1002/zamm.201600215. URL https://asep.lib.cas.cz/ar1-cav/en/detail/?idx=cav_un_epca-1%2A0482472&zf=SHORT.

- [3] Stanislav Sysala, Eva Hruběšová, Zdeněk Michalec, and Franz Tschuchnigg. Optimization and variational principles for the shear strength reduction method. *International Journal for Numerical and Analytical Methods in Geomechanics*, 45(16):2388–2407, 2021. doi: 10.1002/nag.3270. URL <https://tugraz.elsevierpure.com/en/publications/optimization-and-variational-principles-for-the-shear-strength-re-2>.
- [4] Stanislav Sysala, Michal Běreš, Simona Běrešová, Jaroslav Haslinger, Jakub Kružík, and Tomáš Lubér. Convex optimization problems inspired by geotechnical stability analysis. *SIAM Journal on Optimization*, 35(3):1993–2016, 2025. doi: 10.1137/25M1723177. URL <https://doi.org/10.1137/25M1723177>.
- [5] Stanislav Sysala, Michal Běreš, Simona Běrešová, Tomáš Lubér, and Zdeněk Michalec. Advanced continuation and iterative methods for slope stability analysis in 3d. *Computers & Structures*, 315:107842, 2025. doi: 10.1016/j.compstruc.2025.107842. URL <https://doi.org/10.1016/j.compstruc.2025.107842>.
- [6] Anders Logg, Kent-Andre Mardal, and Garth N. Wells, editors. *Automated Solution of Differential Equations by the Finite Element Method: The FEniCS Book*, volume 84 of *Lecture Notes in Computational Science and Engineering*. Springer Berlin Heidelberg, Berlin, Heidelberg, 2012. ISBN 978-3-642-23098-1, 978-3-642-23099-8. doi: 10.1007/978-3-642-23099-8. URL <https://link.springer.com/book/10.1007/978-3-642-23099-8>.
- [7] Igor A. Baratta, Joseph P. Dean, Jørgen S. Dokken, Michal Habera, Jack S. Hale, Chris N. Richardson, Marie E. Rognes, Matthew W. Scroggs, Nathan Sime, and Garth N. Wells. DOLFINx: The next generation FEniCS problem solving environment, 2025. URL <https://zenodo.org/records/18101307>.
- [8] Patrick E. Farrell, David A. Ham, Simon W. Funke, and Marie E. Rognes. Automated derivation of the adjoint of high-level transient finite element programs. *SIAM Journal on Scientific Computing*, 35(4):C369–C393, 2013. doi: 10.1137/120873558. URL <https://doi.org/10.1137/120873558>.
- [9] Sebastian K. Mitusch, Simon W. Funke, and Jørgen S. Dokken. dolfin-adjoint 2018.1: automated adjoints for FEniCS and firedrake. *Journal of Open Source Software*, 4(38):1292, 2019. doi: 10.21105/joss.01292. URL <https://joss.theoj.org/papers/10.21105/joss.01292>.
- [10] Sebastian Blauth. Version 2.0 – cashocs: A computational, adjoint-based shape optimization and optimal control software. *SoftwareX*, 24:101577, 2023. doi: 10.1016/j.softx.2023.101577. URL <https://doi.org/10.1016/j.softx.2023.101577>.
- [11] Tianju Xue, Shuheng Liao, Zhengtao Gan, Chanwook Park, Xiaoyu Xie, Wing Kam Liu, and Jian Cao. JAX-FEM: A differentiable gpu-accelerated 3d finite element solver for automatic inverse design and mechanistic data science. *Computer Physics Communications*, 291:108802, 2023. doi: 10.1016/j.cpc.2023.108802. URL <https://researchportal.hkust.edu.hk/en/publications/jax-fem-a-differentiable-gpu-accelerated-3d-finite-element-solver/>.
- [12] Tianju Xue. Implicit differentiation with second-order derivatives and benchmarks in finite-element-based differentiable physics. *Computer Physics Communications*, 323:110102, 2026. doi: 10.1016/j.cpc.2026.110102. URL <https://doi.org/10.1016/j.cpc.2026.110102>.
- [13] Tobias Bode. AutoPDEx: An automatized partial differential equation solver based on JAX. *Journal of Open Source Software*, 10(108):7300, 2025. doi: 10.21105/joss.07300. URL <https://joss.theoj.org/papers/10.21105/joss.07300>.
- [14] Fanglei Hu, Stephen Niezgoda, Tianju Xue, and Jian Cao. Efficient GPU-computing simulation platform JAX-CPFEM for differentiable crystal plasticity finite element method. *npj Computational Materials*, 11(1):46, 2025. doi: 10.1038/s41524-025-01528-2. URL <https://www.nature.com/articles/s41524-025-01528-2>.
- [15] Ivan Yashchuk. Bringing PDEs to JAX with forward and reverse modes automatic differentiation. *arXiv preprint arXiv:2309.07137*, 2023. doi: 10.48550/arXiv.2309.07137. URL <https://arxiv.org/abs/2309.07137>. arXiv comments state: Published as a workshop paper at ICLR 2020 DeepDiffEq workshop.
- [16] Andrey Latyshev, Jérémy Bleyer, Corrado Maurini, and Jack Hale. Expressing general constitutive models in FEniCSx using external operators and algorithmic automatic differentiation. *Journal of Theoretical, Computational and Applied Mechanics*, 2025. doi: 10.46298/jtcam.14449. URL <https://jtcam.episciences.org/14449>.
- [17] Alberto Cattaneo, M. Keith Ballard, Robert M. Kirby, and Varun Shankar. JetSCI: A hybrid JAX–PETSc framework for scalable differentiable simulation, 2026. URL <https://arxiv.org/abs/2604.22087>.
- [18] Martin Čermák, Stanislav Sysala, and Jan Valdman. Efficient and flexible MATLAB implementation of 2d and 3d elastoplastic problems. *Applied Mathematics and Computation*, 355:595–614, 2019. doi: 10.1016/j.amc.2019.02.054. URL <https://doi.org/10.1016/j.amc.2019.02.054>.
- [19] Ole Sigmund. A 99 line topology optimization code written in MATLAB. *Structural and Multidisciplinary Optimization*, 21(2):120–127, 2001. doi: 10.1007/s001580050176. URL <https://doi.org/10.1007/s001580050176>.

- [20] Martin Philip Bendsøe and Ole Sigmund. *Topology Optimization: Theory, Methods, and Applications*. Springer Berlin Heidelberg, Berlin, Heidelberg, 2 edition, 2004. ISBN 978-3-540-42992-0, 978-3-662-05086-6. doi: 10.1007/978-3-662-05086-6. URL <https://link.springer.com/book/10.1007/978-3-662-05086-6>.
- [21] Federico Ferrari and Ole Sigmund. A new generation 99 line Matlab code for compliance topology optimization and its extension to 3d. *Structural and Multidisciplinary Optimization*, 62(4):2211–2228, 2020. doi: 10.1007/s00158-020-02629-w. URL <https://orbit.dtu.dk/en/publications/a-new-generation-99-line-matlab-code-for-compliance-topology-opti/>.
- [22] Blaise Bourdin. Filters in topology optimization. *International Journal for Numerical Methods in Engineering*, 50(9):2143–2158, 2001. doi: 10.1002/nme.116. URL <https://doi.org/10.1002/nme.116>.
- [23] Yingqi Jia, Chao Wang, and Xiaojia Shelly Zhang. FEniTop: a simple FEniCSx implementation for 2d and 3d topology optimization supporting parallel computing. *Structural and Multidisciplinary Optimization*, 67(8):140, 2024. doi: 10.1007/s00158-024-03818-7. URL <https://link.springer.com/article/10.1007/s00158-024-03818-7>.
- [24] Atilim Gunes Baydin, Barak A. Pearlmutter, Alexey Andreyevich Radul, and Jeffrey Mark Siskind. Automatic differentiation in machine learning: a survey. *Journal of Machine Learning Research*, 18(153):1–43, 2018. URL <https://www.jmlr.org/papers/v18/17-468.html>.
- [25] Roy Frostig, Matthew James Johnson, and Chris Leary. Compiling machine learning programs via high-level tracing. In *SysML*, 2018. URL <https://research.google/pubs/compiling-machine-learning-programs-via-high-level-tracing/>.
- [26] Juan C. Simo and Robert L. Taylor. Consistent tangent operators for rate-independent elastoplasticity. *Computer Methods in Applied Mechanics and Engineering*, 48(1):101–118, 1985. doi: 10.1016/0045-7825(85)90070-2. URL [https://doi.org/10.1016/0045-7825\(85\)90070-2](https://doi.org/10.1016/0045-7825(85)90070-2).
- [27] Juan C. Simo and Thomas J. R. Hughes. *Computational Inelasticity*, volume 7 of *Interdisciplinary Applied Mathematics*. Springer-Verlag, New York, 1998. ISBN 978-0-387-97520-7, 978-0-387-22763-4. doi: 10.1007/b98904. URL <https://link.springer.com/book/10.1007/b98904>.
- [28] A. R. Curtis, M. J. D. Powell, and J. K. Reid. On the estimation of sparse jacobian matrices. *Journal of the Institute of Mathematics and its Applications*, 13(1):117–119, 1974. doi: 10.1093/imamat/13.1.117. URL <https://academic.oup.com/imamat/article-abstract/13/1/117/721987>.
- [29] Thomas F. Coleman and Jorge J. Moré. Estimation of sparse jacobian matrices and graph coloring problems. *SIAM Journal on Numerical Analysis*, 20(1):187–209, 1983. doi: 10.1137/0720013. URL <https://doi.org/10.1137/0720013>.
- [30] Thomas F. Coleman and Jorge J. Moré. Estimation of sparse hessian matrices and graph coloring problems. *Mathematical Programming*, 28(3):243–270, 1984. doi: 10.1007/BF02612334. URL <https://link.springer.com/article/10.1007/BF02612334>.
- [31] Satish Balay, William D. Gropp, Lois Curfman McInnes, and Barry F. Smith. Efficient management of parallelism in object-oriented numerical software libraries. In Erlend Arge, Are Magnus Bruaset, and Hans Petter Langtangen, editors, *Modern Software Tools for Scientific Computing*, pages 163–202. Birkhäuser Boston, Boston, MA, 1997. ISBN 978-0-8176-3974-7, 978-1-4612-1986-6. doi: 10.1007/978-1-4612-1986-6_8. URL https://link.springer.com/chapter/10.1007/978-1-4612-1986-6_8.
- [32] Satish Balay, Shrirang Abhyankar, Mark F. Adams, Steven Benson, Jed Brown, Peter Brune, Kris Buschelman, Emil M. Constantinescu, Lisandro Dalcin, Alp Dener, Victor Eijkhout, Jacob Faibussowitsch, William D. Gropp, Václav Hapla, Tobin Isaac, Pierre Jolivet, Dmitry Karpeev, Dinesh Kaushik, Matthew G. Knepley, Fande Kong, Scott Kruger, Dave A. May, Lois Curfman McInnes, Richard Tran Mills, Lawrence Mitchell, Todd Munson, Jose E. Roman, Karl Rupp, Patrick Sanan, Jason Sarich, Barry F. Smith, Stefano Zampini, Hong Zhang, Hong Zhang, and Junchao Zhang. PETSc web page. <https://petsc.org/>, 2026. URL <https://petsc.org/>. Official PETSc 3.25 citation block, accessed 2026-04-30.
- [33] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*. Springer Series in Operations Research and Financial Engineering. Springer New York, New York, NY, 2 edition, 2006. ISBN 978-0-387-30303-1, 978-0-387-40065-5. doi: 10.1007/978-0-387-40065-5. URL <https://link.springer.com/book/10.1007/978-0-387-40065-5>.
- [34] Peter Lindqvist. *Notes on the Stationary p-Laplace Equation*. SpringerBriefs in Mathematics. Springer International Publishing, Cham, 2019. ISBN 978-3-030-14500-2, 978-3-030-14501-9. doi: 10.1007/978-3-030-14501-9. URL <https://link.springer.com/book/10.1007/978-3-030-14501-9>.
- [35] V. L. Ginzburg and L. D. Landau. On the theory of superconductivity. *Zhurnal Eksperimental’noi i Teoreticheskoi Fiziki*, 20(12):1064–1082, 1950. URL <https://cds.cern.ch/record/486430>.

- [36] Javier Bonet and Richard D. Wood. *Nonlinear Continuum Mechanics for Finite Element Analysis*. Cambridge University Press, Cambridge, 2 edition, 2008. ISBN 978-0-521-83870-2, 978-0-511-75544-6. doi: 10.1017/CBO9780511755446. URL <https://doi.org/10.1017/CBO9780511755446>.
- [37] E. H. Davis. Theories of plasticity and the failure of soil masses. In I. K. Lee, editor, *Soil Mechanics: Selected Topics*, pages 341–380. Butterworths, London, 1968. URL <https://trid.trb.org/View/124976>.
- [38] IT4Innovations National Supercomputing Center. Karolina: Compute nodes, 2026. URL <https://docs.it4i.cz/en/docs/clusters/karolina/compute-nodes>. Accessed 2026-05-05.